

Lecture #09 - Parallel Computing and OpenMP

AMath 483/583

Announcements

- Homework #1 Solutions posted within a week
- Homework #2 Updates —> add remote, pull (demo)
- fixes will be made this morning, announced

Word of Wisdom

- For single-core and parallel computing...

“Premature optimization is the root of all evil.” —
Donald Knuth

(Computer scientist, mathematician, “Father of algorithm analysis.”)

Parallel Computing

- Old Version:

processor speed
doubles every
eighteen months

- New Version:

number of cores
doubles every
eighteen months



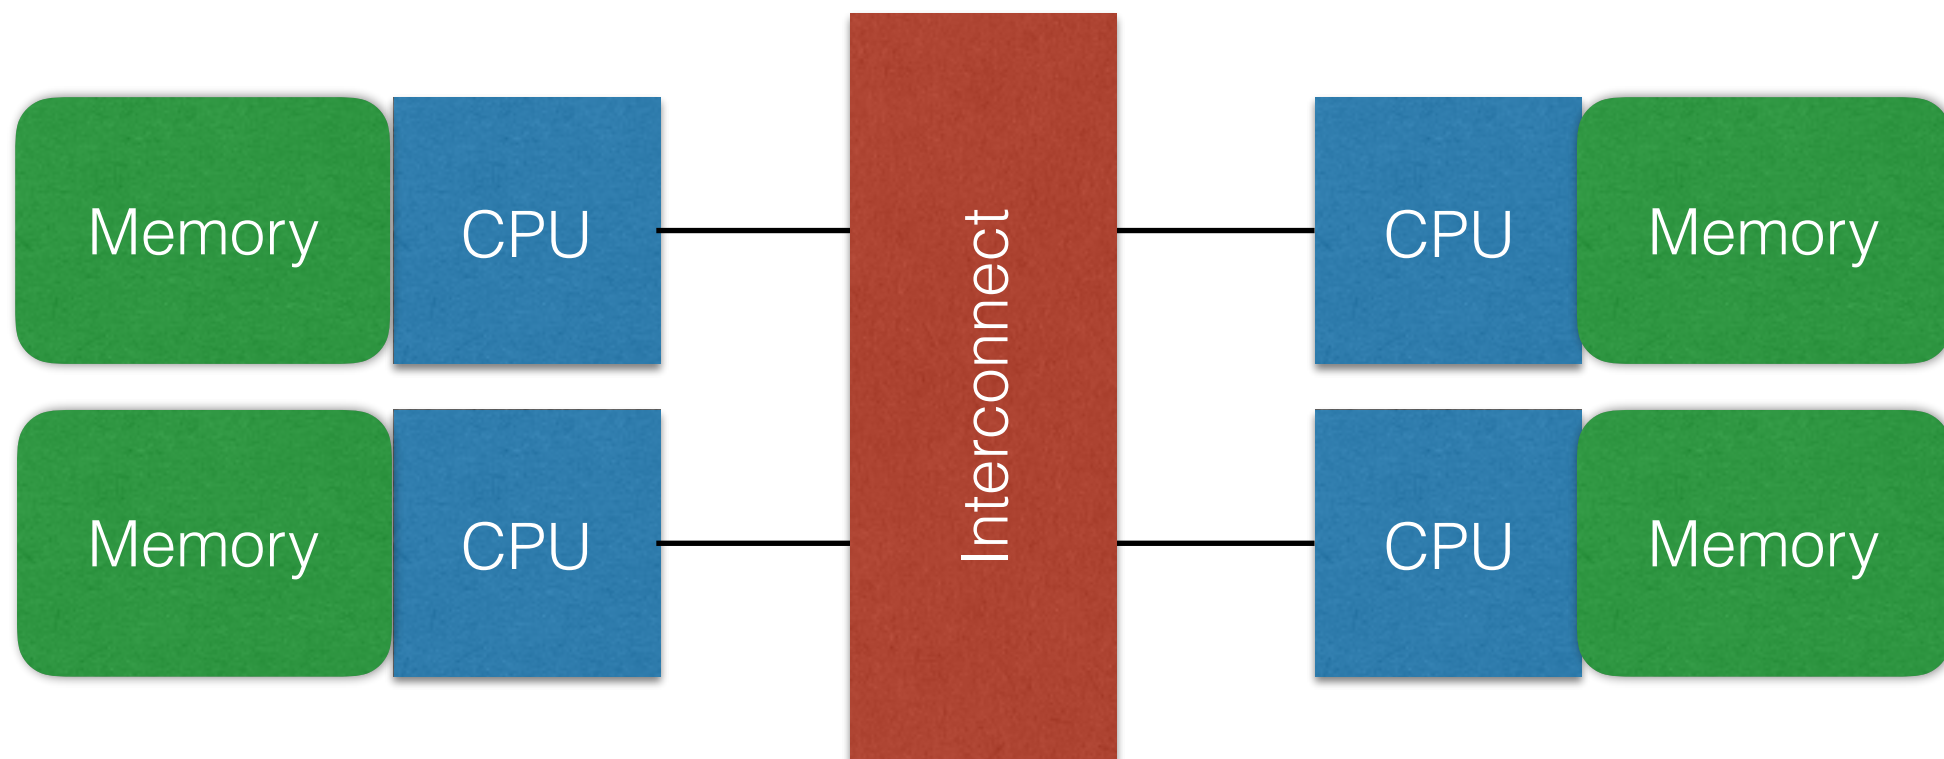
Shared Memory

- All processors / nodes have access to same memory
 - e.g. L3 cache (in most chips) and RAM
 - nodes on massively parallel machines
 - implicit data communication



Distributed Memory

- Each processor / node has its own memory pool
 - non-example: L1 and L2 cache on most processors (*coherence*)
 - nodes on massively parallel machines
 - explicit data communication



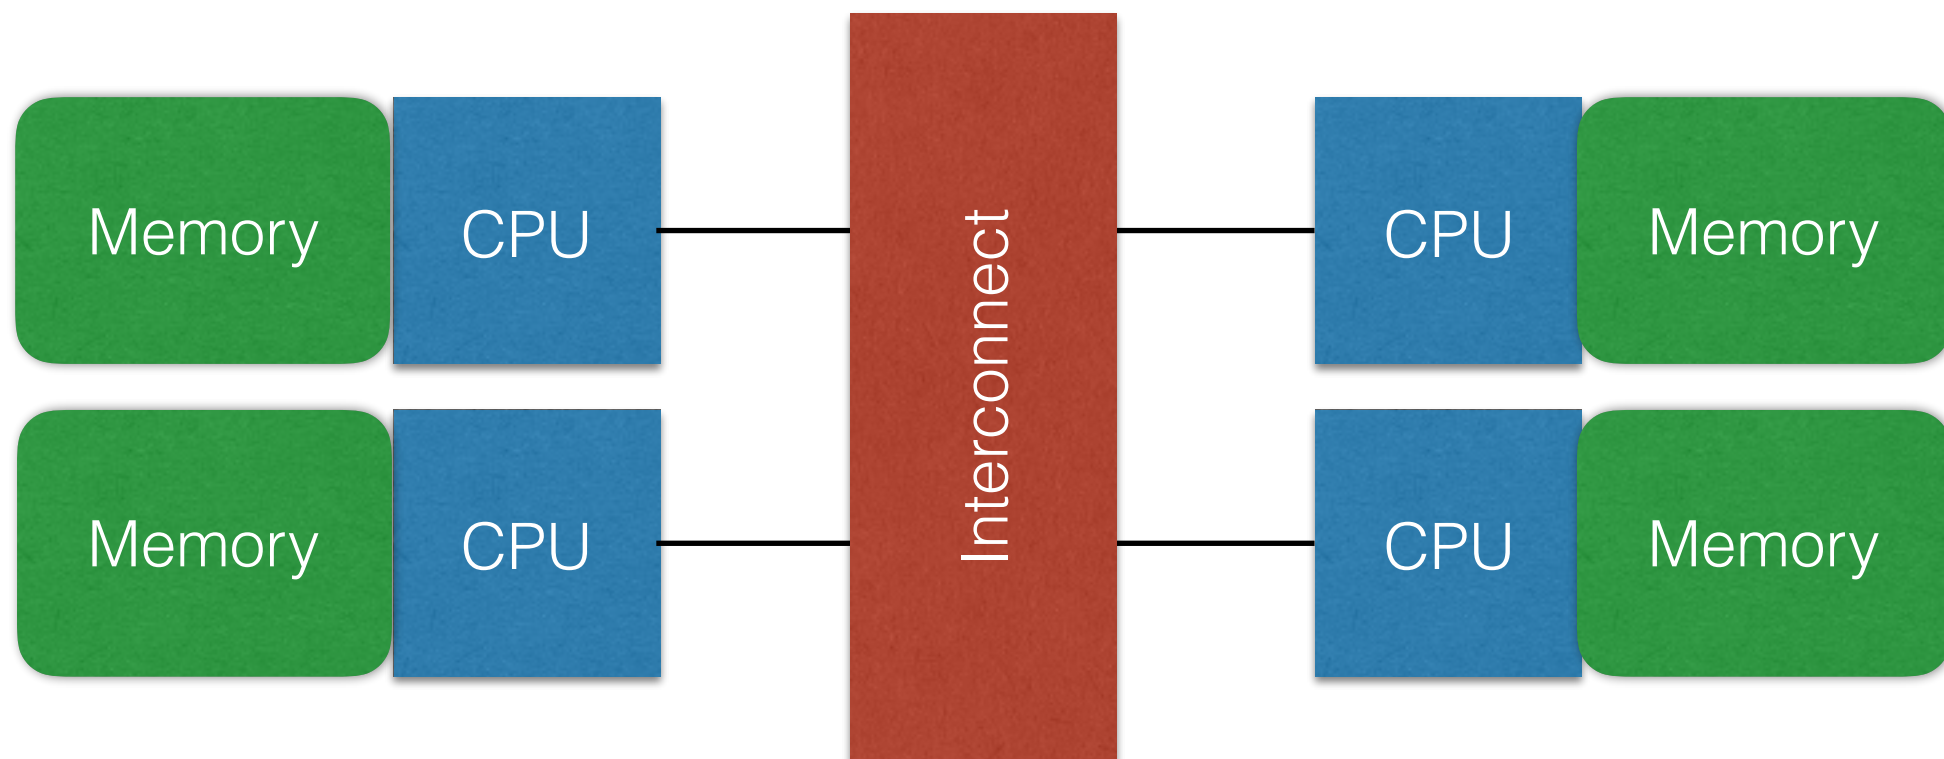
Distributed Memory

- Each processor
 - non-exclusive access to local memory
 - nodes communicate via a shared interconnect
 - explicit data communication

The interconnect is very slow!
(Relative to CPU \longleftrightarrow Shmem communication)

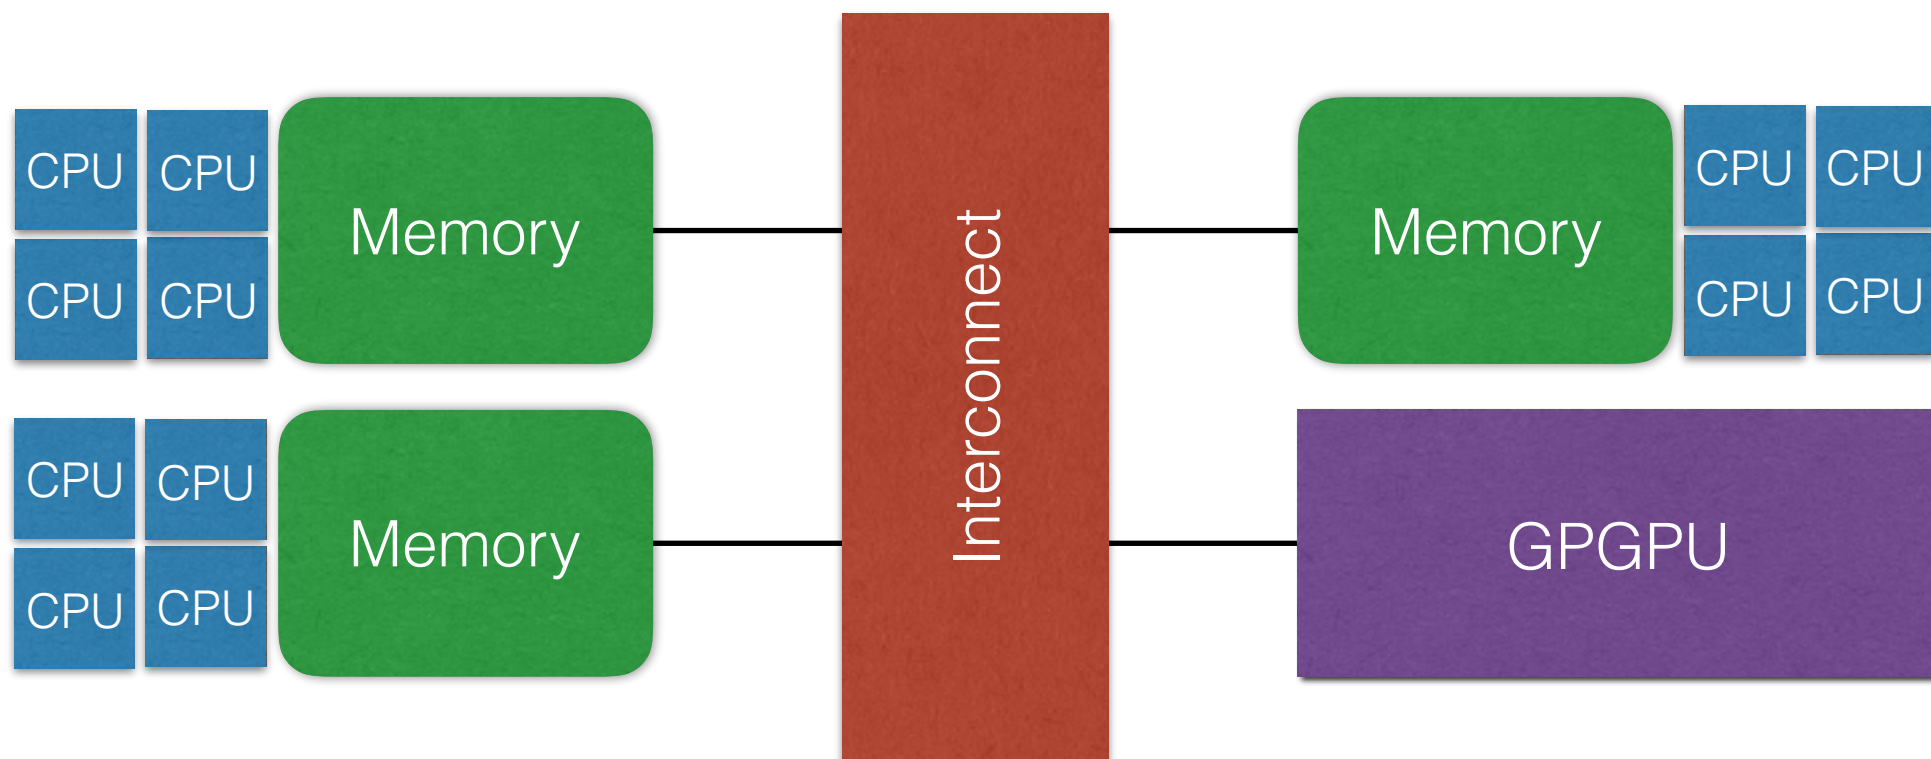
CPU \longrightarrow RAM = 100 ns

CPU \longrightarrow CPU = 10,000+ ns



Multi-Hardware Systems

- Combined systems of shared + distributed memory
- Dedicated hardware:
 - SIMD / Vector Processors
 - GPGPUs

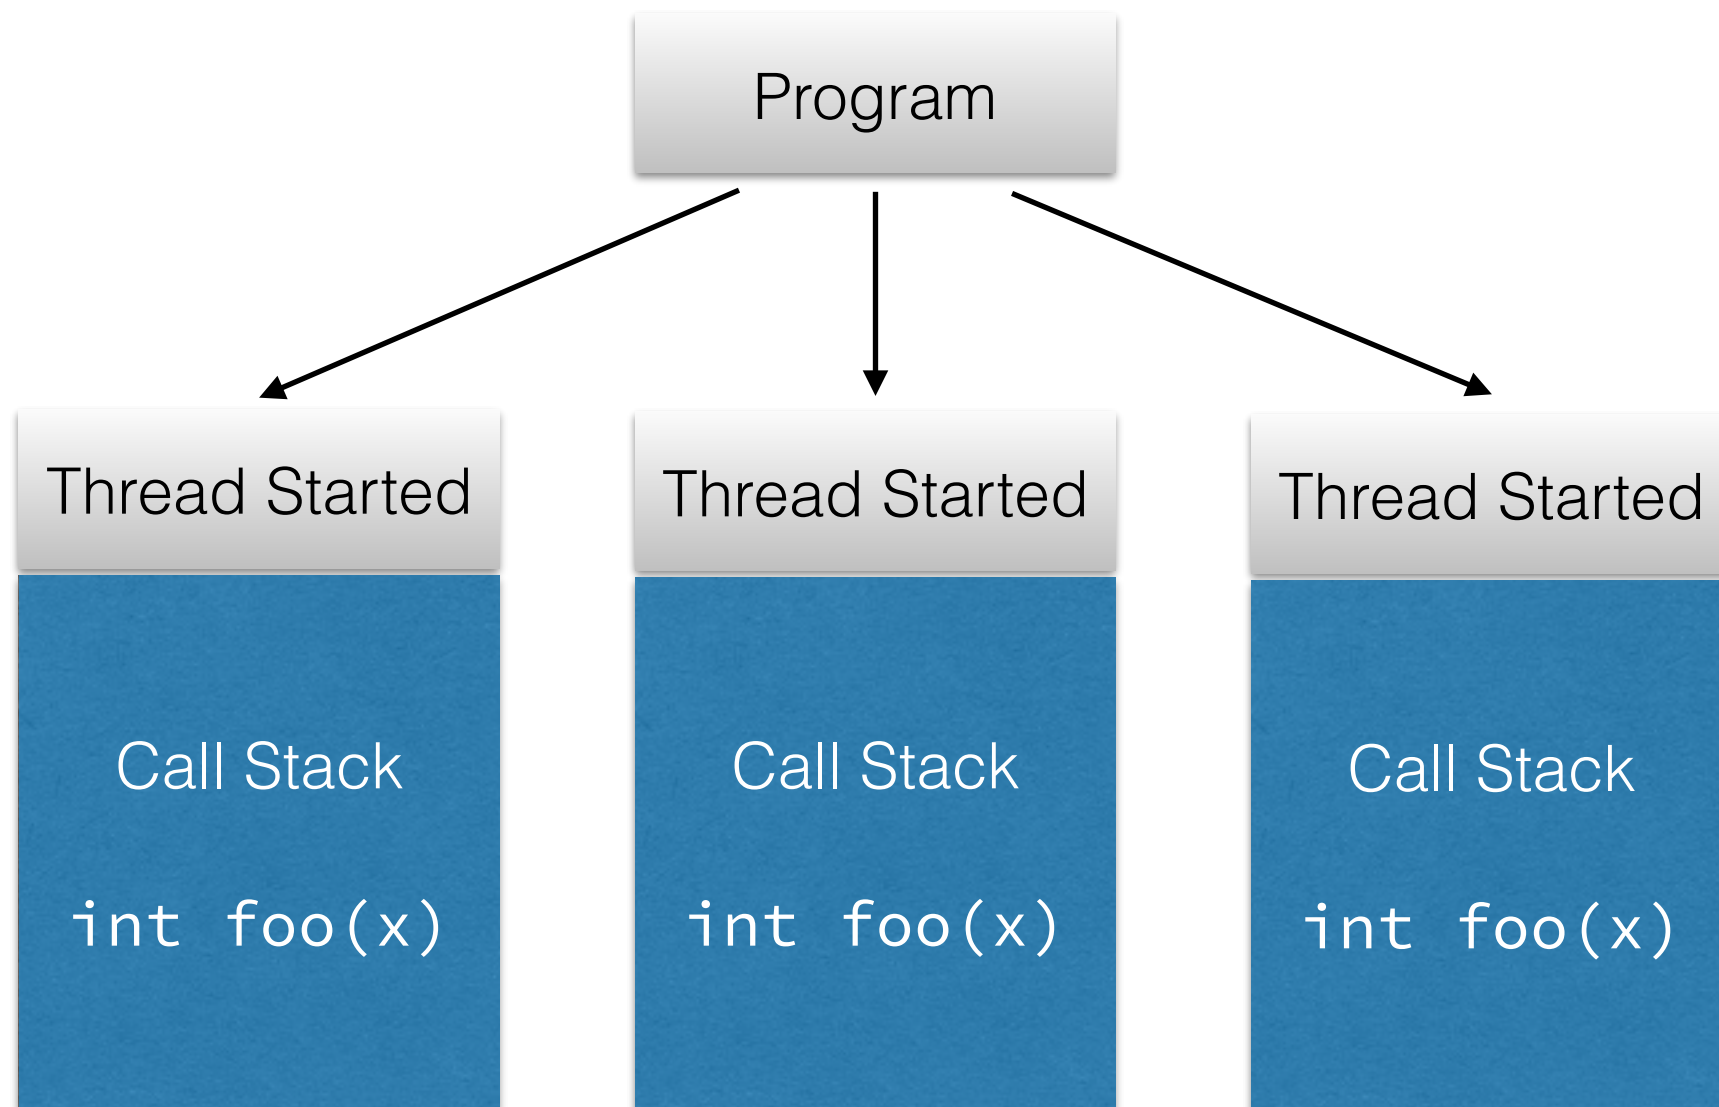


Threads

- *“Sequence of instructions with both thread-specific data and access to shared address space.”*
- each thread has own call stack and local vars
- each thread can access shared data
- Spawned and destroyed throughout a program / computation.

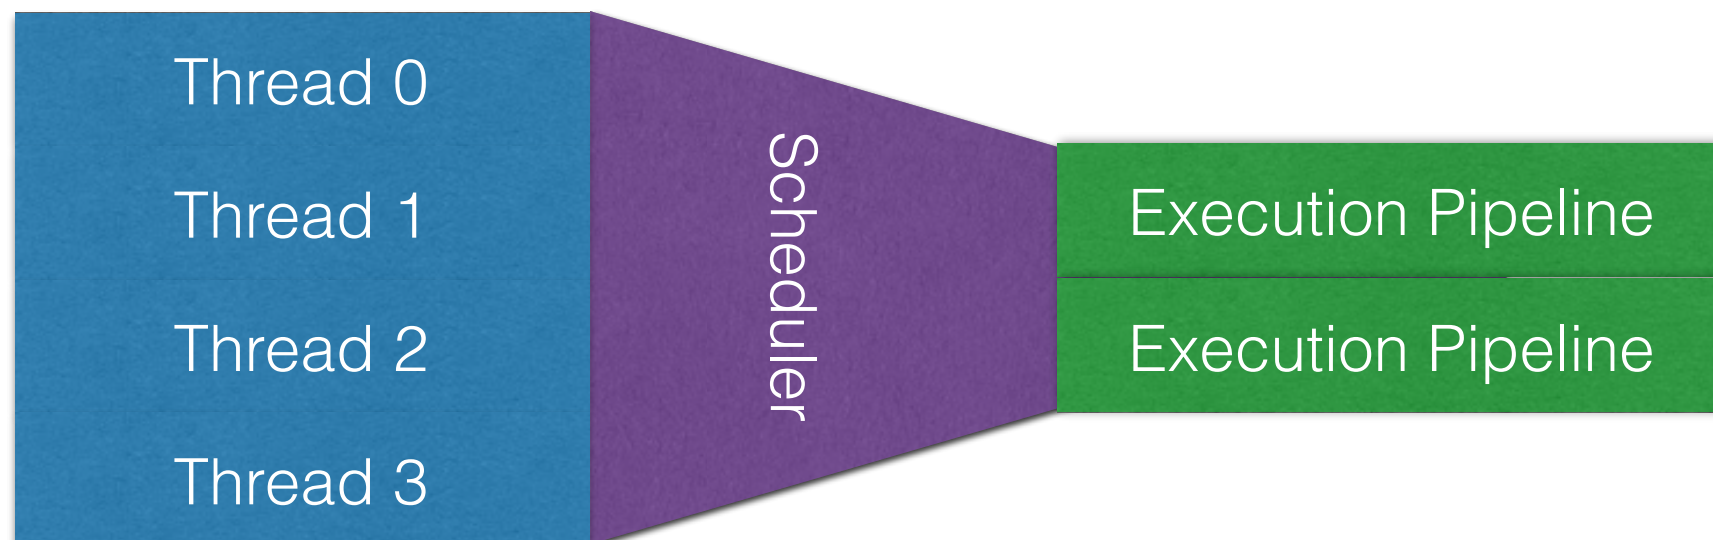
Threads

- Each thread has its own call stack (even when running same code)



Threads

- Software Threads — spawned by program
- Hardware Threads — simult. instruction pipelines offered by CPU
 - “scheduler” feeds software threads into hardware
 - schedule chooses which instruction to execute next (major area of research)



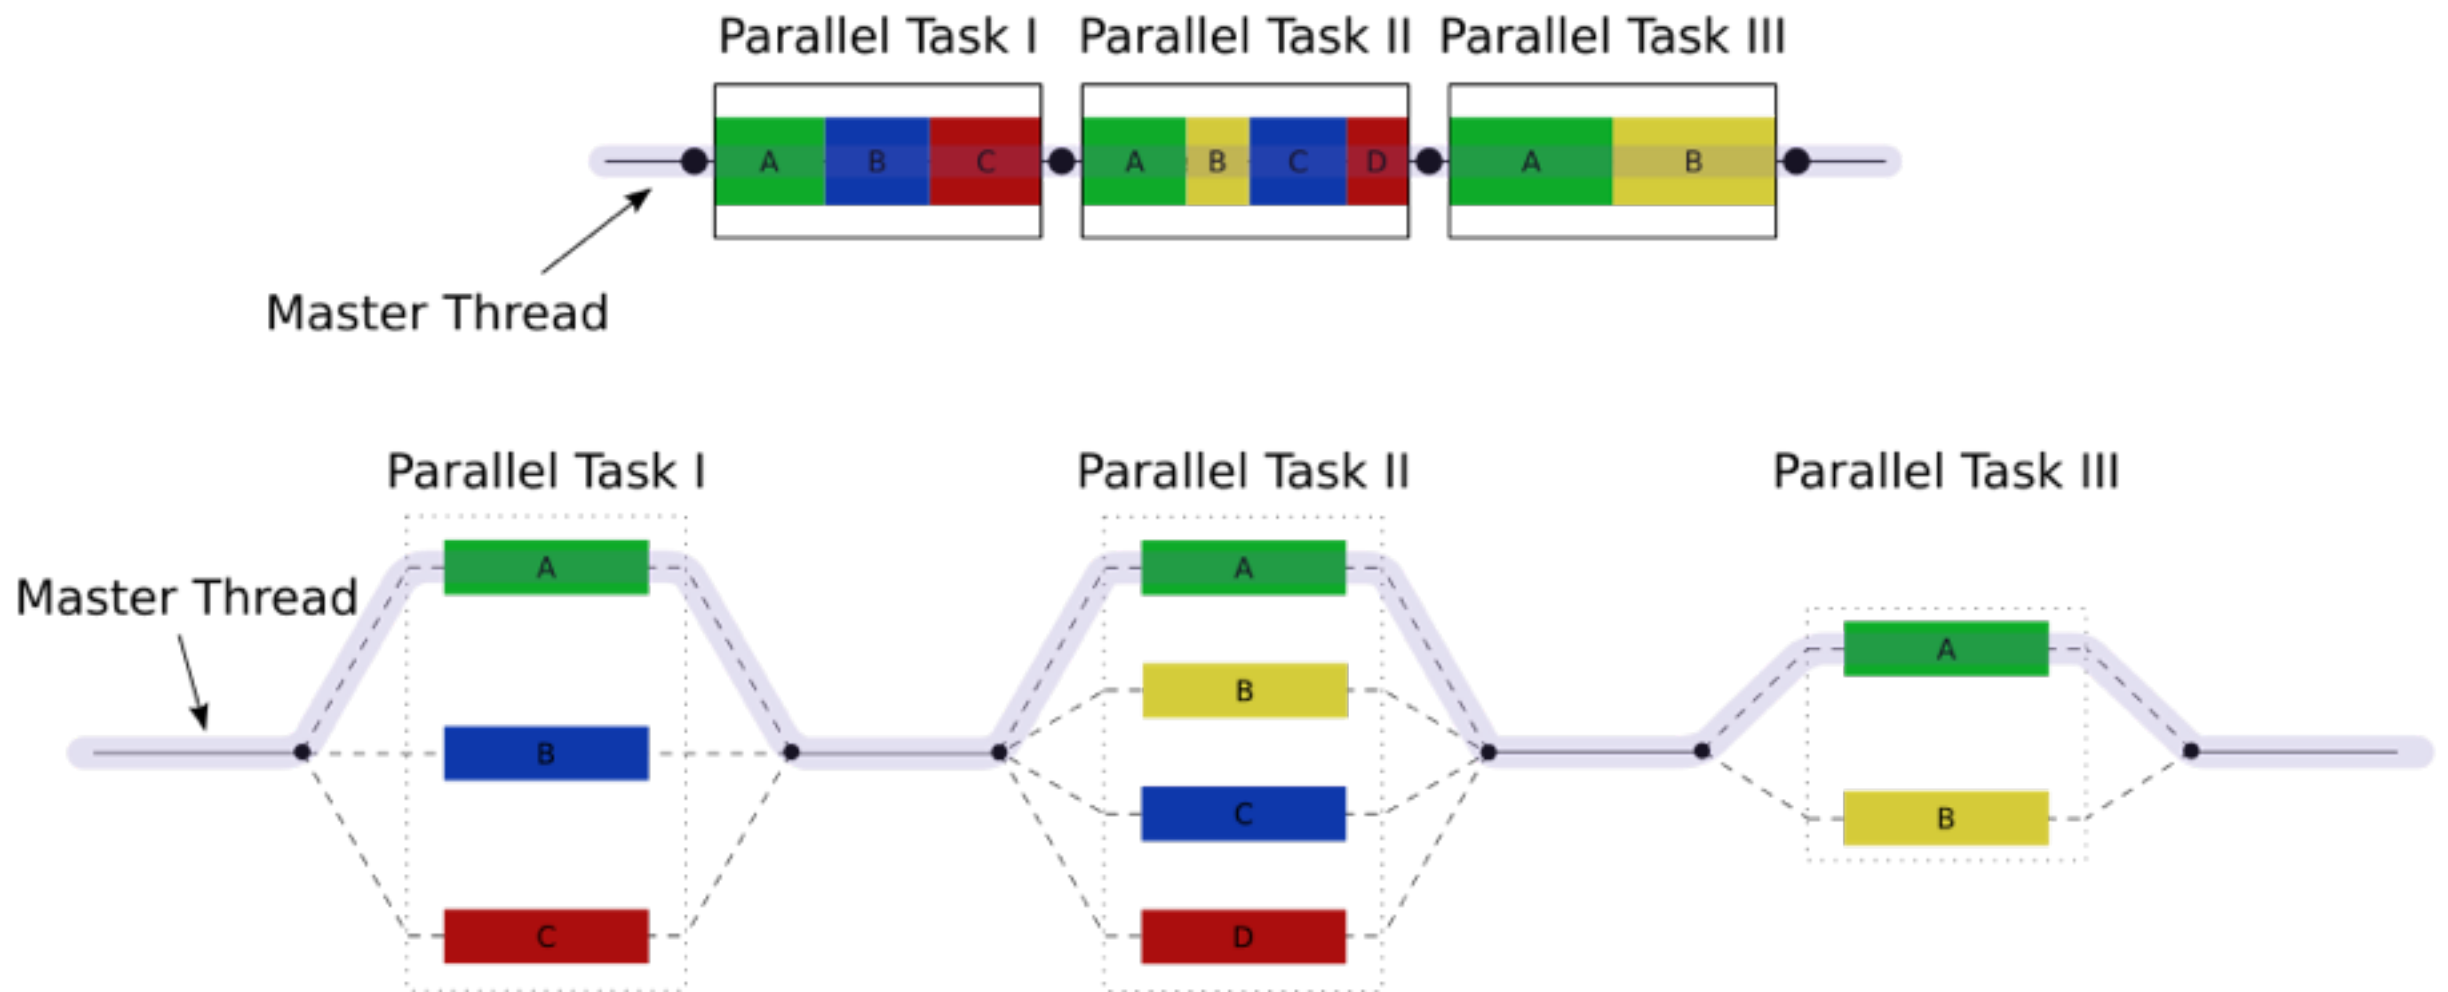
Parallel Work

- “[Embarrassingly parallel \(link\)](#)” — each thread works completely independently
- `vec_add` — each thread adds across a component: $\text{out}[i] = v[i] + w[i]$
- `mat_mul` — each thread computes C_{ij}
- `gradient_descent` — each thread tries a different initial guess x_0

Parallel Work

- Some situations more complicated:
 - `vec_norm` — how do you parallelize summation across N elements?
 - Problem: multiple threads writing to same result
- Need to break parallelism to “*synchronize*” threads.

Parallel Work

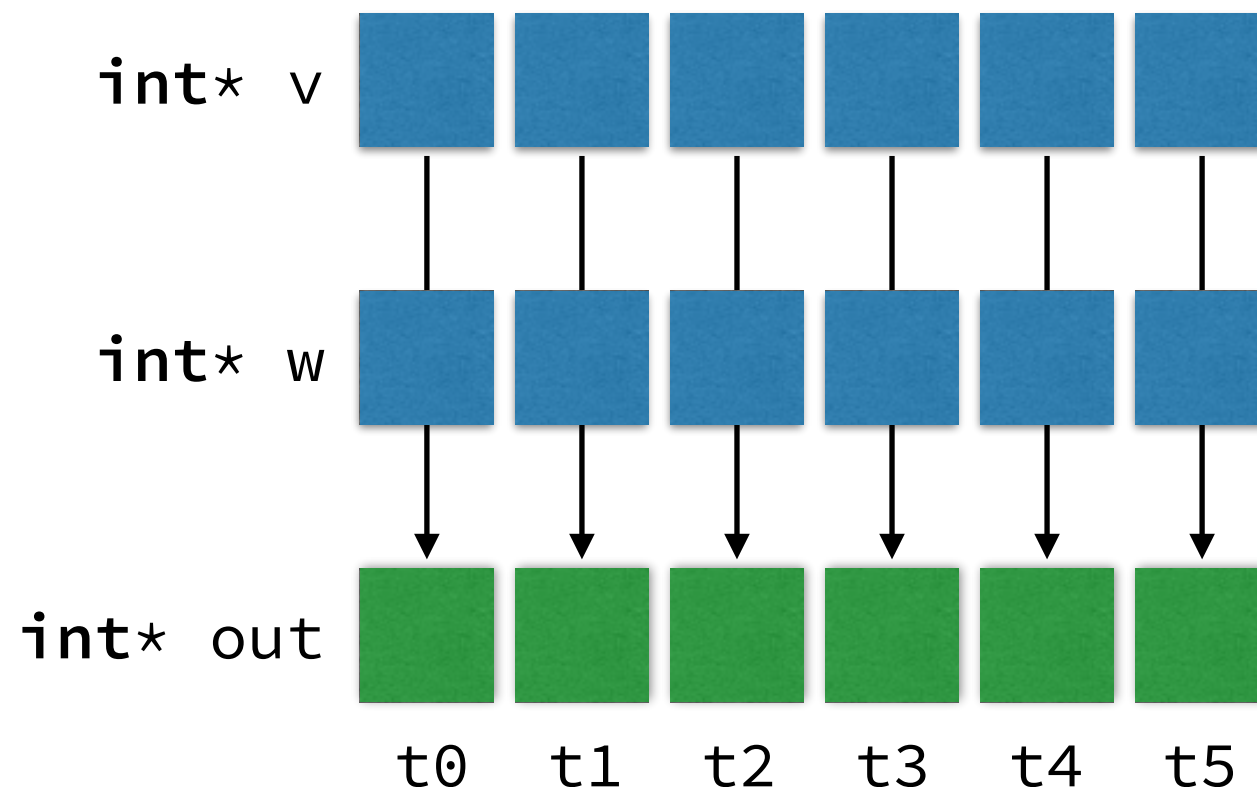


OpenMP

- Shared memory multiprocessing in C/C++/FORTRAN
 - easy to use (difficult to debug)
 - wraps PThreads
- [POSIX Threads](#) (link)
 - standardized C thread programming library
 - UNIX standard since 1995

Threads Issues

- Multiple threads have access to same data — very deep subject
- **Simple Situation:** `vec_add` — each thread computes the sum at each index, data is independent / disjoint



Threads Issues

- **Complex Situation:** count — determine the number of elements in an array greater than x

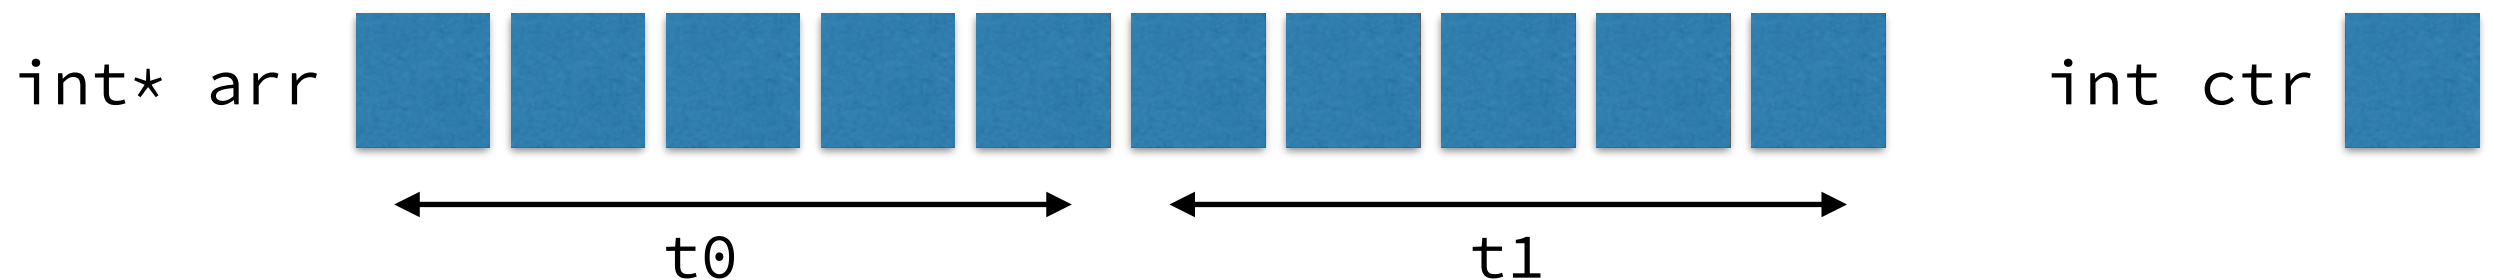


```
for (int i=0; i<length; ++i)
    if (arr[i] < x)
        ctr += 1;
```

- Three step process:
 - read counter, add one, write to counter

Threads Issues

- **Complex Situation:** count — determine the number of elements in an array greater than x



```
// Executed by thread 0
for (int i=0; i<length/2; ++i)
{
    if (arr[i] < x)
        ctr += 1;
}
// Executed by thread 1
for (int i=length/2; i<length; ++i)
{
    if (arr[i] < x)
        ctr += 1;
}
```

Thread 0: read counter,
increment, write to counter

Thread 1: read counter,
increment, write to counter

Threads Issues

- Disjoint Threads

int ctr



time



Thread 0:



Thread 1:

read
increment
write

Threads Issues

- Overlapping Threads

int ctr



time



Thread 0:



Thread 1:



Threads Issues

- Overlapping Threads

Problem: Multiple Threads, Same Data



Must Synchronize Actions!
(Mutex / condition variables / blocks.)

Thread 0

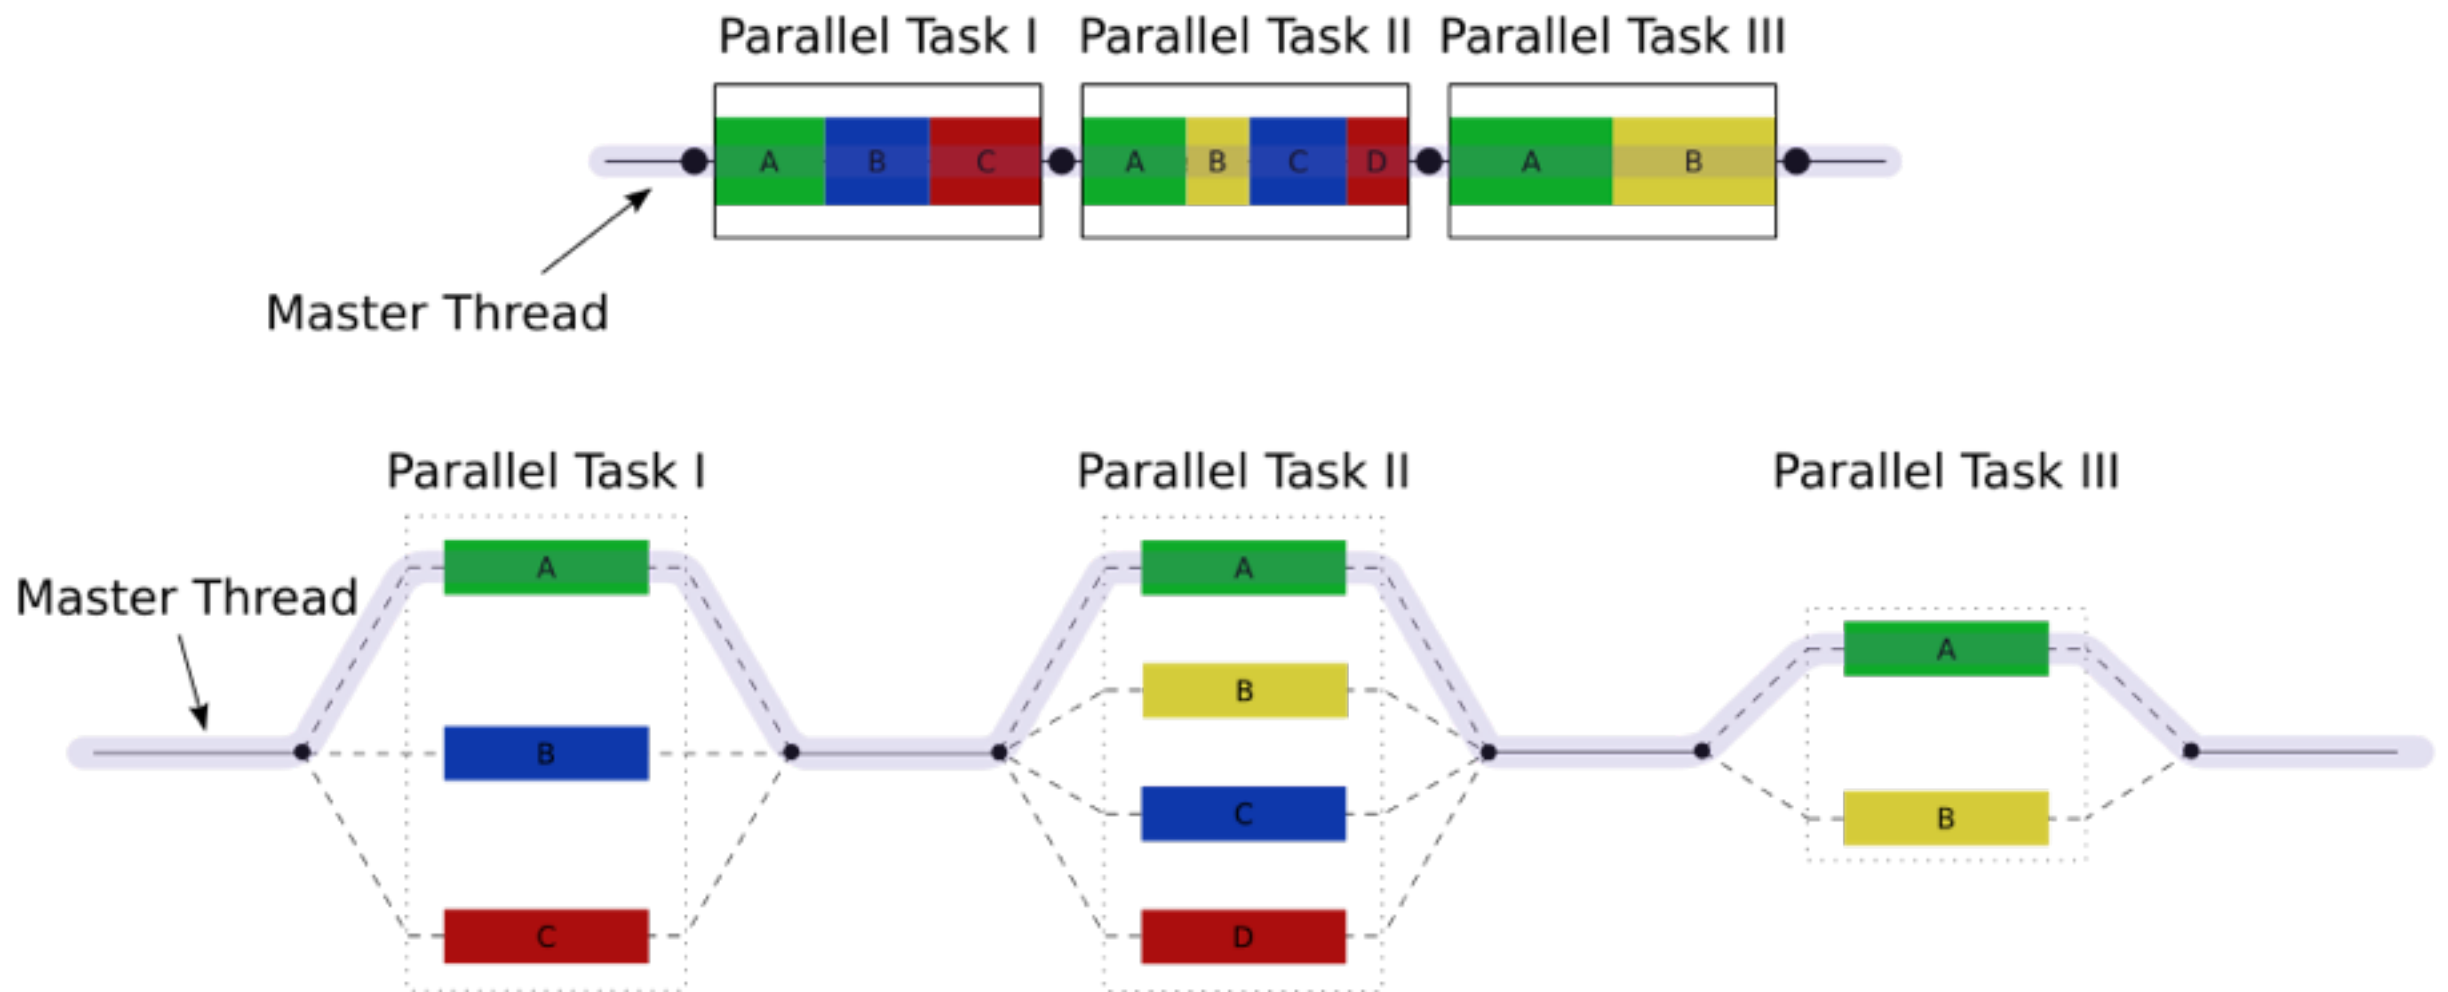
Thread 1:

read

increment

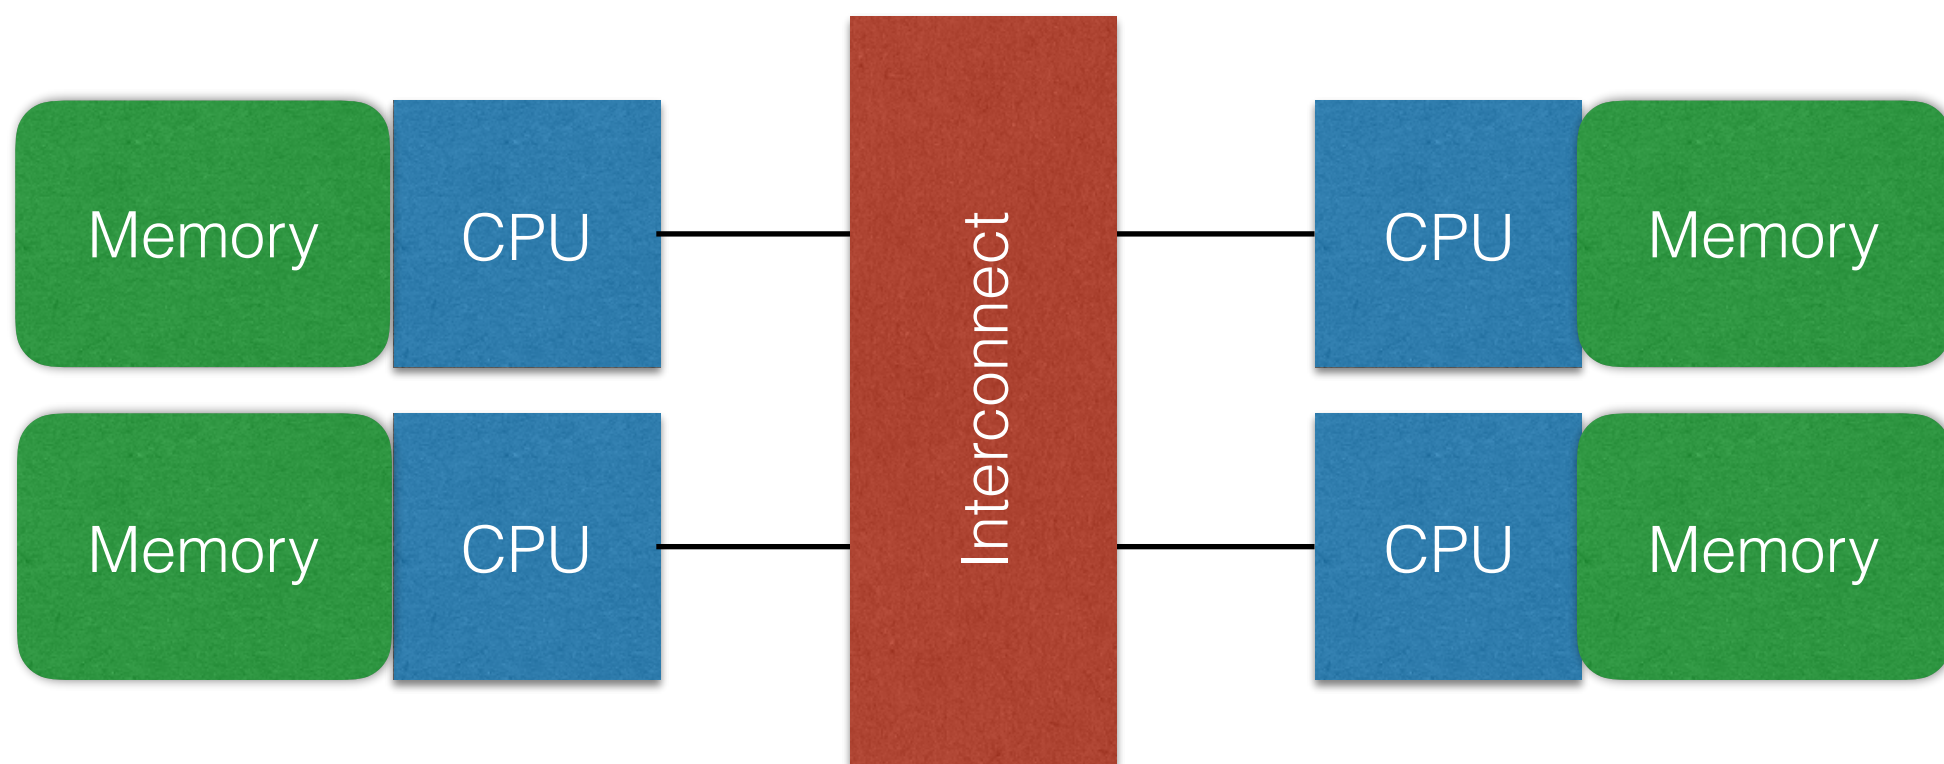
write

Thread Issues



Processes

- “Thread with its own address space.” (i.e. thread’s data is private)
- Thousands+ processes running on computer
 - (I can give this talk, move my mouse cursor, receive (many) emails...)
- Process model used in distributed memory systems:



MPI

- “Message Passing Interface” for C/C++/FORTRAN
- Processes communicate by explicitly passing messages to each other
 - high communication overhead —> maximize local data processing before sending info
- “Coarse-grain parallelism”
- Uses sockets / TCP over network

Amdahl's Law

- Only part of a computation can be parallelized

Suppose 50% of computation is inherently sequential and other 50% is parallelizable.

- **Question:** given N processors how much faster can this computation be run?

Amdahl's Law

- Only part of a computation can be parallelized

Suppose 50% of computation is inherently sequential and other 50% is parallelizable.

- **Question:** given N processors how much faster can this computation be run?
- **Answer:** at most x2 faster (*indep. of N !*)
 - (if parallel part reduced to zero time)

Amdahl's Law

- Only part of a computation can be parallelized

*Suppose **10%** of computation is inherently sequential and other **90%** is parallelizable.*

- **Question:** given N processors how much faster can this computation be run?

Amdahl's Law

- Only part of a computation can be parallelized

*Suppose **10%** of computation is inherently sequential and other **90%** is parallelizable.*

- **Question:** given N processors how much faster can this computation be run?
- **Answer:** at most $\times 10$ faster (*indep. of N !*)
 - (sequential part is still taking $1/10$ of time)

Amdahl's Law

Suppose $1/S$ of the computation is inherently sequential and other $(1 - 1/S)$ can be parallelized.

- At most factor of S speedup
- Let T_S be time required in sequential (single-core) machine to run a computation

$$T_S = (1/S)T_S + (1 - 1/S)T_S$$

Amdahl's Law

- Now run on P processors
- Let T_P be time required on parallel machine. Then T_P is at least:

$$T_P = (1/S)T_S + \frac{(1-1/S)}{P}T_S$$

- Note:

$$\lim_{P \rightarrow \infty} T_P = (1/S)T_S$$

Amdahl's Law

- *Suppose $1/S$ of the computation is inherently sequential, then*

$$T_P = (1/S)T_S + \frac{(1-1/S)}{P}T_S$$

where T_S is the original sequential time and T_P is the time taken across P processors

Speedup

- **Speedup** = T_S / T_P
 - Typically, speedup $\ll P$
- Amdahl's law does not account for overhead costs:
 - starting / destroying processes and threads, thread communication, ...

$$T_P = (1/S)T_S + (1 - 1/S)T_S + T_{\text{overhead}}$$

Scaling

- Some algorithms scale better than others as P increases...
 - embarrassingly parallelizable algorithms
 - algorithms with low / batched communication between threads
 - few blocking calls
- Let N = problem size (solve $N \times N$ linear system, simulate N particles in space, etc.)

Strong Scaling

- *“How does the algorithm perform as the number of processors P increases for **fixed problem size N** ?”*

Any algorithm will eventually break down. (Consider $P > N$.)

Weak Scaling

- *“How does the algorithm perform when the problem size increases **with the number of processors?**”*

e.g. will doubling the number of processors allow us to solve a problem twice as large in the same time?

Weak Scaling

- “*Twice as large*”
- **Example:** Solving $N \times N$ linear system with Gaussian elimination (LU-factorization) = $\mathcal{O}(N^3)$
- double the problem size = increase N by a factor of $2^{1/3} \approx 1.26$
- e.g. $100 \times 100 \longrightarrow 126 \times 126$

Weak Scaling

- “Twice

- Example
Gauss

- double
of 2

- e.g.

Warning!

$\mathcal{O}(N^3)$ refers to number of multiplications.

Memory accesses are much more expensive. How does that change with N ?

with

$$\mathcal{O}(N^3)$$

the N by a factor

Lecture #10 - OpenMP

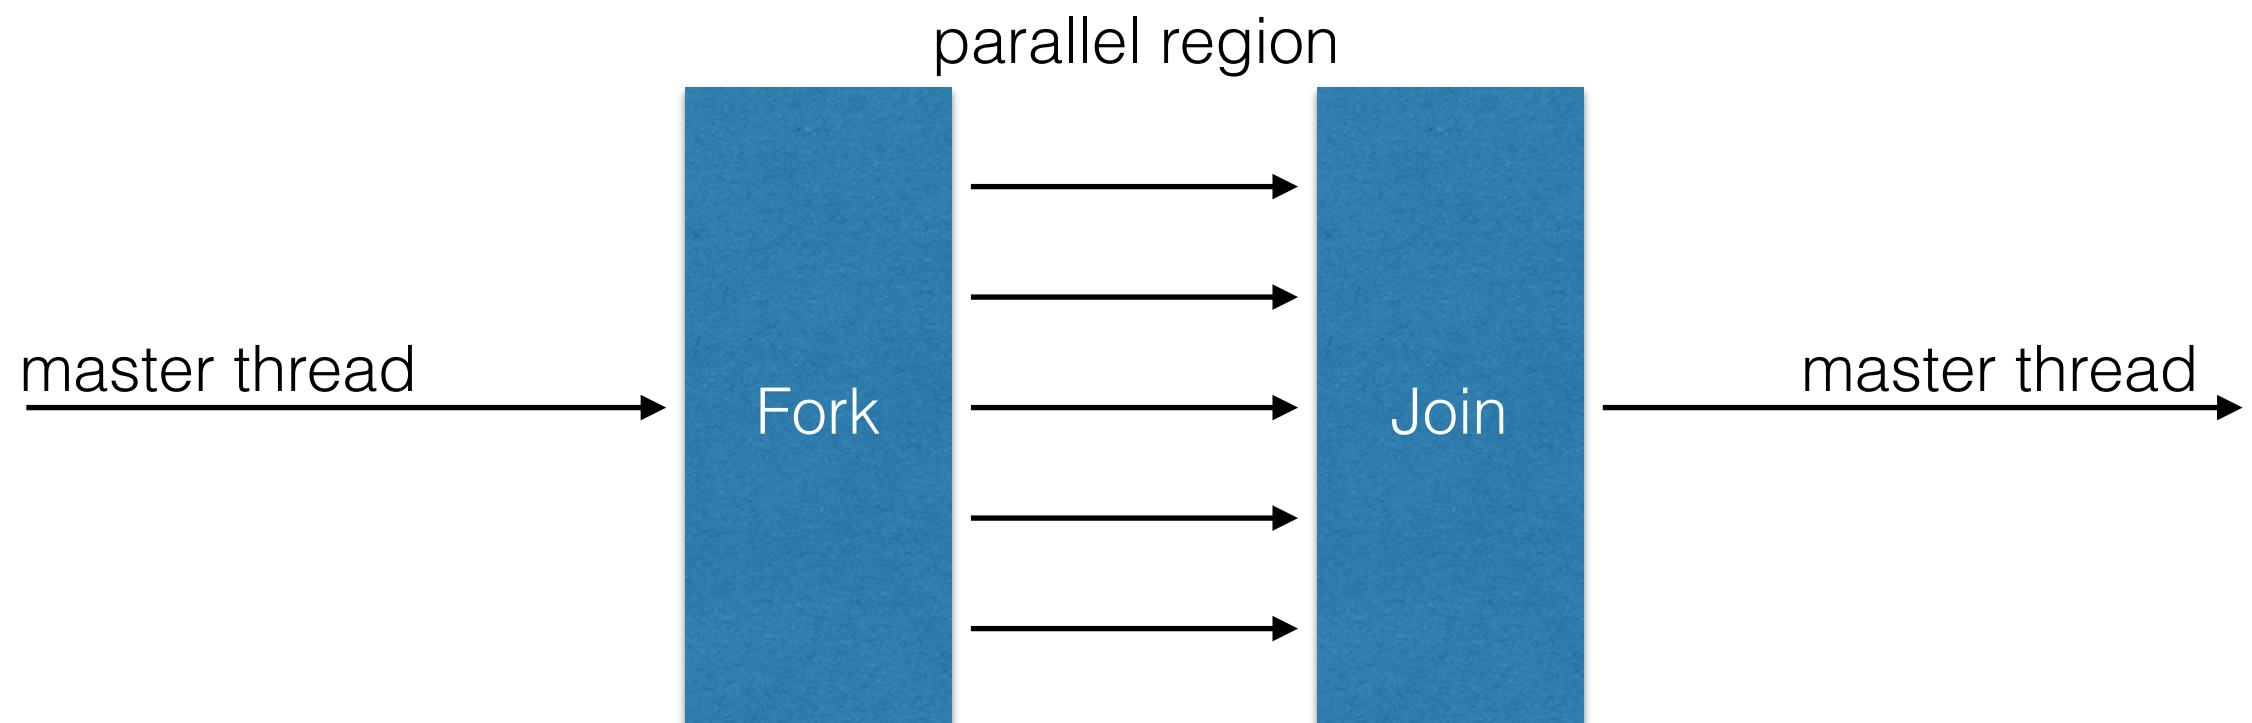
AMath 483/583

OpenMP Resources

- <http://www.openmp.org/wp/resources>
- Tim Mattson's (Intel) "Introduction to OpenMP" (YouTube video series)
- Lawrence Livermore National Labs "Introduction to OpenMP" (Tutorial + Reference)
- + many more
- *(theme of this class: lectures are a starting point, online resources are abundant!)*

OpenMP - Basic Idea

- “Fork - Join” model
 - begin with single thread (master thread)
 - FORK: master thread creates team of parallel threads
 - JOIN: when threads complete action they synchronize and terminate



OpenMP - Basic Idea

- Compiler directives — how to fork / join
 - how to assign *parallel regions* to threads
 - what data is *private* (local) to each thread
- Compiler **generates** multithreaded code
- Dependencies:
 - remove them OR explicitly synchronize

OpenMP Compiler Directives

- Within C code:

```
#pragma omp [directivename] [options]
```

- Compile C code containing OpenMP directives:

```
$ gcc -fopenmp ...
```

OpenMP Compiler Directives

```
#pragma omp [directive] [clause ...]  
    if (scalar_expression)  
    private (list)  
    shared (list)  
    default (shared || none)  
    firstprivate (list)  
    reduction (operator: list)  
    copyin (list)  
    num_threads (integer-expression)
```

OpenMP Compiler Directives

- Examples

```
#pragma omp parallel [clause]
{
    // block of parallel code
}
```

```
#pragma omp parallel for [clause]
for (...)
{
    // for loop body
}
```

```
#pragma omp barrier
// wait for all threads to arrive here before proceeding
```

Hello World

- Each thread prints “hello, world”

```
#include “omp.h”
```

OpenMP header file

```
int main()
```

```
{
```

Parallel region with default number of threads

```
#pragma omp parallel
```

```
{
```

Runtime library function to get current thread #

```
    int id = omp_get_thread_num();
```

```
    printf(“Hello(%d),”, id);
```

```
    printf(“world(%d)”, id);
```

```
}
```

End of parallel region

```
}
```

Demo

OpenMP Hello World