

Lecture #10 - OpenMP

AMath 483/583

Announcements

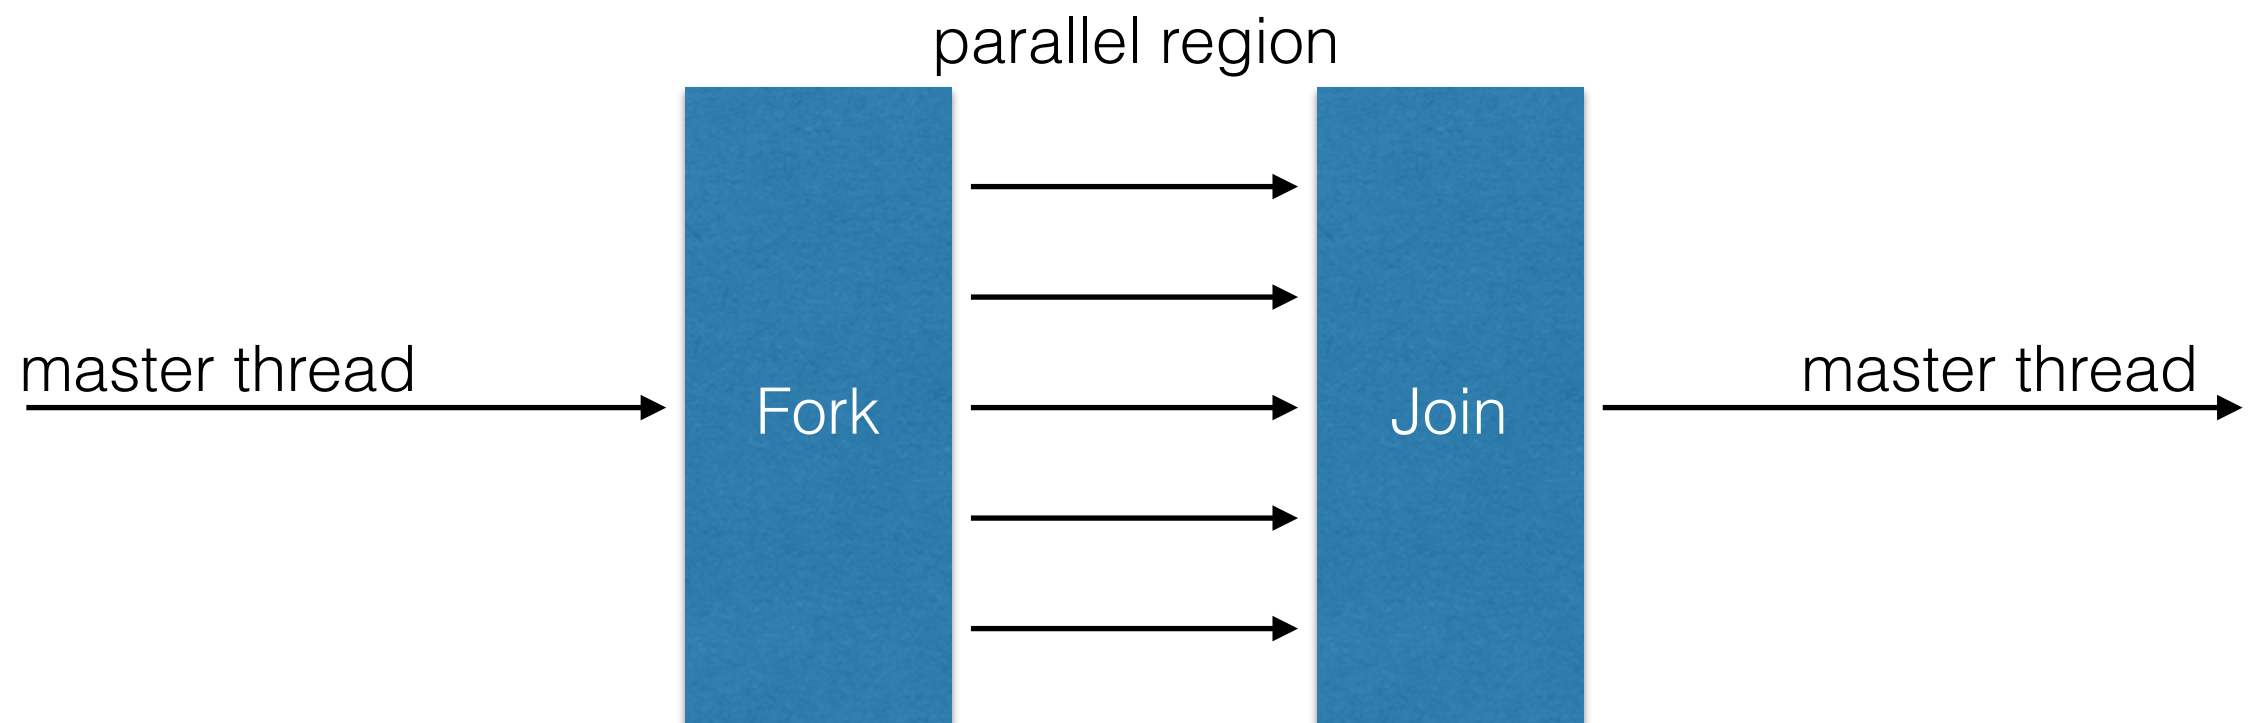
- Remember: Homework #2 due next Friday, not Monday
- References and code will be available later today

Note

- SMC currently will only give you two cores
- Four later
- Today's demo is on my four-core machine
 - warning default compiler on OS X is “clang”
 - “gcc-5”

OpenMP - Basic Idea

- “Fork - Join” model
 - begin with single thread (master thread)
 - FORK: master thread creates team of parallel threads
 - JOIN: when threads complete action they synchronize and terminate



OpenMP Compiler Directives

```
#pragma omp [directive] [clause ...]  
    if (scalar_expression)  
    private (list)  
    shared (list)  
    default (shared || none)  
    firstprivate (list)  
    reduction (operator: list)  
    copyin (list)  
    num_threads (integer-expression)
```

OpenMP Compiler Directives

- Examples

```
#pragma omp parallel [clause]
{
    // block of parallel code
}
```

```
#pragma omp parallel for [clause]
for (...)
{
    // for loop body
}
```

```
#pragma omp barrier
// wait for all threads to arrive here before proceeding
```

Hello World

- Each thread prints “hello, world”

```
#include “omp.h”
```

OpenMP header file

```
int main()
```

```
{
```

Parallel region with default number of threads

```
#pragma omp parallel
```

```
{
```

Runtime library function to get current thread #

```
    int id = omp_get_thread_num();
```

```
    printf(“Hello(%d),”, id);
```

```
    printf(“world(%d)”, id);
```

```
}
```

End of parallel region

```
}
```

Hello World

- Each thread prints "hello, world"

```
#include
```

```
int
```

```
{
```

```
#pragma omp
```

```
{
```

```
int id = omp_get_thread_num();
```

```
printf("Hello(%d),", id);
```

```
printf("world(%d)", id);
```

```
}
```

```
}
```

Compile Using

```
$ gcc -fopenmp hello_openmp.c -o hello
```

rary function to
ent thread #

End of parallel region

Setting Thread Numbers

- Statically (at compile-time)

```
#pragma omp parallel thread_num(4)
{
    // parallel code
}
```

- Dynamically (at run-time)

```
omp_set_num_threads(4); // a run-time function
#pragma omp parallel
{
    // parallel code
}
```



Demo

Setting run-time threads via command line:
`dynamic_threads.c`

Private Variables

- Variables declared *inside* a parallel region are private to each thread

```
#pragma omp parallel
{
    int id = omp_get_thread_num();
    ...
}
```




Private Variables

- Variables declared *outside* a parallel region are shared / accessible by each thread

```
int id;  
pragma omp parallel  
{  
    id = omp_get_thread_num();  
    ...  
}
```


Danger! Multiple threads are writing to same location.



Private Variables

- **Solution:** variables can be declared “private” —> each thread will have a local copy

```
int id;  
pragma omp parallel private(id)  
{  
    id = omp_get_thread_num();  
    ...  
}
```



Demo

Private variables

Loops

- Share iterations of a loop across threads
- useful in “data parallel” situations

```
int i;  
for (i=0; i<10; ++i)  
    // loop body
```

Loops

- **Default:** each thread gets one iteration
get next available iteration when complete

```
int i;  
#pragma omp parallel private(i) num_threads(4)  
{  
    #pragma omp for  
    for (i=0; i<10; ++i)  
    {  
        // loop body  
    }  
}
```


Loops

- **Default:** each thread gets one iteration
get next available iteration when complete

```
int i;  
#pragma omp parallel private(i) num_threads(4)  
{  
    #pragma omp for  
    for (i=0; i<10; ++i)  
    {  
        // loop body  
    }  
}
```

Declare parallel
region

Loops

- **Default:** each thread gets one iteration
get next available iteration when complete

```
int i;  
#pragma omp parallel private(i) num_threads(4)  
{  
    #pragma omp for  
    for (i=0; i<10; ++i)  
    {  
        // loop body  
    }  
}
```

Tell threads to split up work
in for loop

Loops

- **Default:** each thread gets one iteration
get next available iteration when complete

```
int i;  
#pragma omp parallel private(i) num_threads(4)  
{  
    for (i=0; i<10; ++i)  
    {  
        // loop body  
    }  
}
```

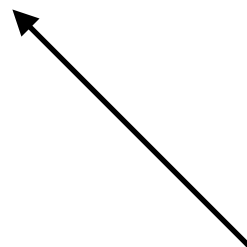
WARNING

If “#pragma omp for” is omitted then all threads do same for loop

Loops

- **Example:** `vec_add`

```
void vec_add(double* out, double* v, double* w, double* N)
{
    #pragma omp parallel shared(out,v,w) num_threads(4)
    {
        int i; // implicitly private
        #pragma omp for
        for (i=0; i<N; ++i)
            out[i] = v[i] + w[i]
    }
}
```



Although implicit, good idea to explicitly declare “shared” for readability.

Demo

Parallel vec_add and timing

Shortcut

- Combine parallel directives and for directives

```
#pragma omp parallel for [num_threads(N)]  
for (. . .)  
    { . . . }
```

- Convenience function.

Reminder: Private Variables

- Consider the following:

```
double x, dx = 1.0 / (N + 1.0);
```

```
for (int i=0; i<N; ++i)  
{  
    x = i*dx;  
    y[i] = sin(x) * cos(x);  
}
```

Reminder: Private Variables

- **Incorrect:**

```
double x, dx = 1.0 / (N + 1.0);
```

```
#pragma omp parallel for  
for (int i=0; i<N; ++i)  
{  
    x = i*dx;  
    y[i] = sin(x) * cos(x);  
}
```


Reminder: Private Variables

- **Incorrect:**

```
double x, dx = 1.0 / (N + 1.0);
```

```
#pragma omp parallel for  
for (int i=0; i<N; ++i)  
{  
    x = i*dx;  
    y[i] = sin(x) * cos(x);  
}
```

By default, x is a shared variable.

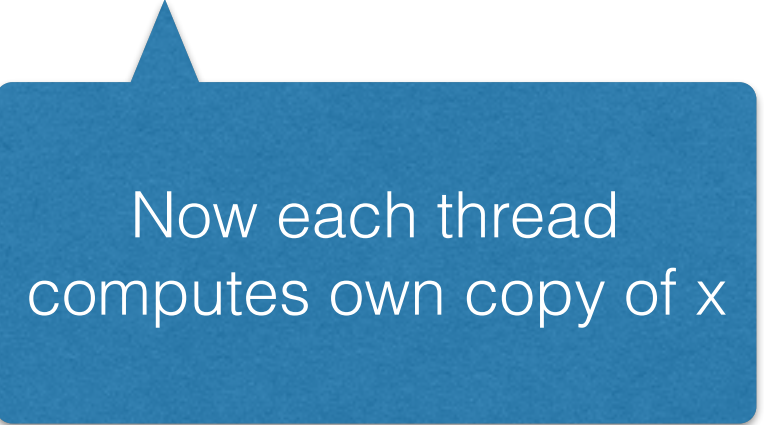
Thread 0: set x
Thread 1: set different x
Thread 0: evaluate y[i]

Reminder: Private Variables

- **Incorrect:**

```
double x, dx = 1.0 / (N + 1.0);
```

```
#pragma omp parallel for private(x)
for (int i=0; i<N; ++i)
{
    x = i*dx;
    y[i] = sin(x) * cos(x);
}
```



Now each thread
computes own copy of x

Loop Chunking

- `vec_add` — work done by each thread is small

Thread 0

`out[0] = v[0] + w[0]; out[5] = v[5] + w[5]; ...`

Thread 1

`out[2] = v[2] + w[2]; out[4] = v[4] + w[4]; ...`

...

- Thread scheduler works hard at finding next iteration for each thread

Loop Chunking

- Instead, give each thread a “chunk” of iterations

Thread 0

do iterations $i=0$ through $i=(M-1)$

Thread 1

do iterations $i=M$ through $i=(2*M-1)$

...

- In some instances this is more efficient
- *Static vs. dynamic* chunking - assign at start or as you go?

Loop Chunking

- **schedule:** declare chunking behavior

```
void vec_add(double* out, double* v, double* w, double* N)
{

    #pragma omp parallel shared(out,v,w) num_threads(4)
    {
        int i;
        #pragma omp for
        for (i=0; i<N; ++i)
            out[i] = v[i] + w[i]
    }
}
```

Loop Chunking

- **schedule:** declare chunking behavior

```
void vec_add(double* out, double* v, double* w, double* N)
{
    #pragma omp parallel shared(out,v,w) num_threads(4)
    {
        int i;
        #pragma omp for schedule(static,128)
        for (i=0; i<N; ++i)
            out[i] = v[i] + w[i]
    }
}
```

Set chunk size to 128 and use static scheduling

(each thread gets 128 iterations of the loop at a time)

Loop Chunking

- **schedule:** declare chunk size

Chunk size can be stored in
(shared) variable.

(Each thread needs to know.)

```
void vec_add(double* out, double* v, double* w, double* N)
{
    int chunk_size = 128;
    #pragma omp parallel shared(out,v,w,chunk_size) num_threads(4)
    {
        int i;
        #pragma omp for schedule(static,chunk_size)
        for (i=0; i<N; ++i)
            out[i] = v[i] + w[i]
    }
}
```

Chunking Strategies

- `schedule(static [,chunk])` — deal out blocks of iterations of size “chunk” to each thread
- `schedule(dynamic [,chunk])` — each thread grabs “chunk” iterations off of queue until all iterations have been handled
- `schedule(guided [,chunk])` — threads dynamically grab blocks of iterations. Block size starts large and shrinks down to “chunk” size
- `schedule(auto)` — runtime can “learn” from previous executions of same loop

Demo

Chunking and performance tuning

Synchronization

- Impose order constraints and protect access to shared data
- Case study:

```
#pragma omp parallel
{
    int id = omp_get_thread_num();
    A[id] = // some big calculation

    // use A in some function to compute Bi
    B[id] = func(A, id);
}
```

Synchronization

- Impose order constraints and protect access to shared data
- Case study:

```
#pragma omp parallel
```

```
{
```

```
    int id = omp_get_thread_rank();
```

```
    A[id] = // some big calculation
```

```
    // use A in some function
```

```
    B[id] = func(A, id);
```

```
}
```

Problem!

A is not necessarily fully
formed at this point!

(Need to wait for all threads)

Synchronization - Barrier

- Impose order constraints and protect access to shared data
- Case study:

```
#pragma omp parallel
```

```
{
```

```
    int id = omp_get_thread_num();
```

```
    A[id] = // some big calculation
```

```
    #pragma omp barrier
```

Threads wait here until all threads arrive

```
    // use A in some function to compute Bi
```

```
    B[id] = func(A, id);
```

```
}
```

Synchronization - Barrier

- Some OpenMP directives have natural barriers

```
#pragma omp for
```

```
for (. . .)
```

```
    { . . . }
```

```
// all threads synchronize at end of loop
```

```
// before proceeding
```

```
#pragma omp for nowait
```

```
for (. . .)
```

```
    { . . . }
```

```
// thread i will not wait for thread j to
```

```
// finish at last iterations of loop
```

Synchronization - Critical

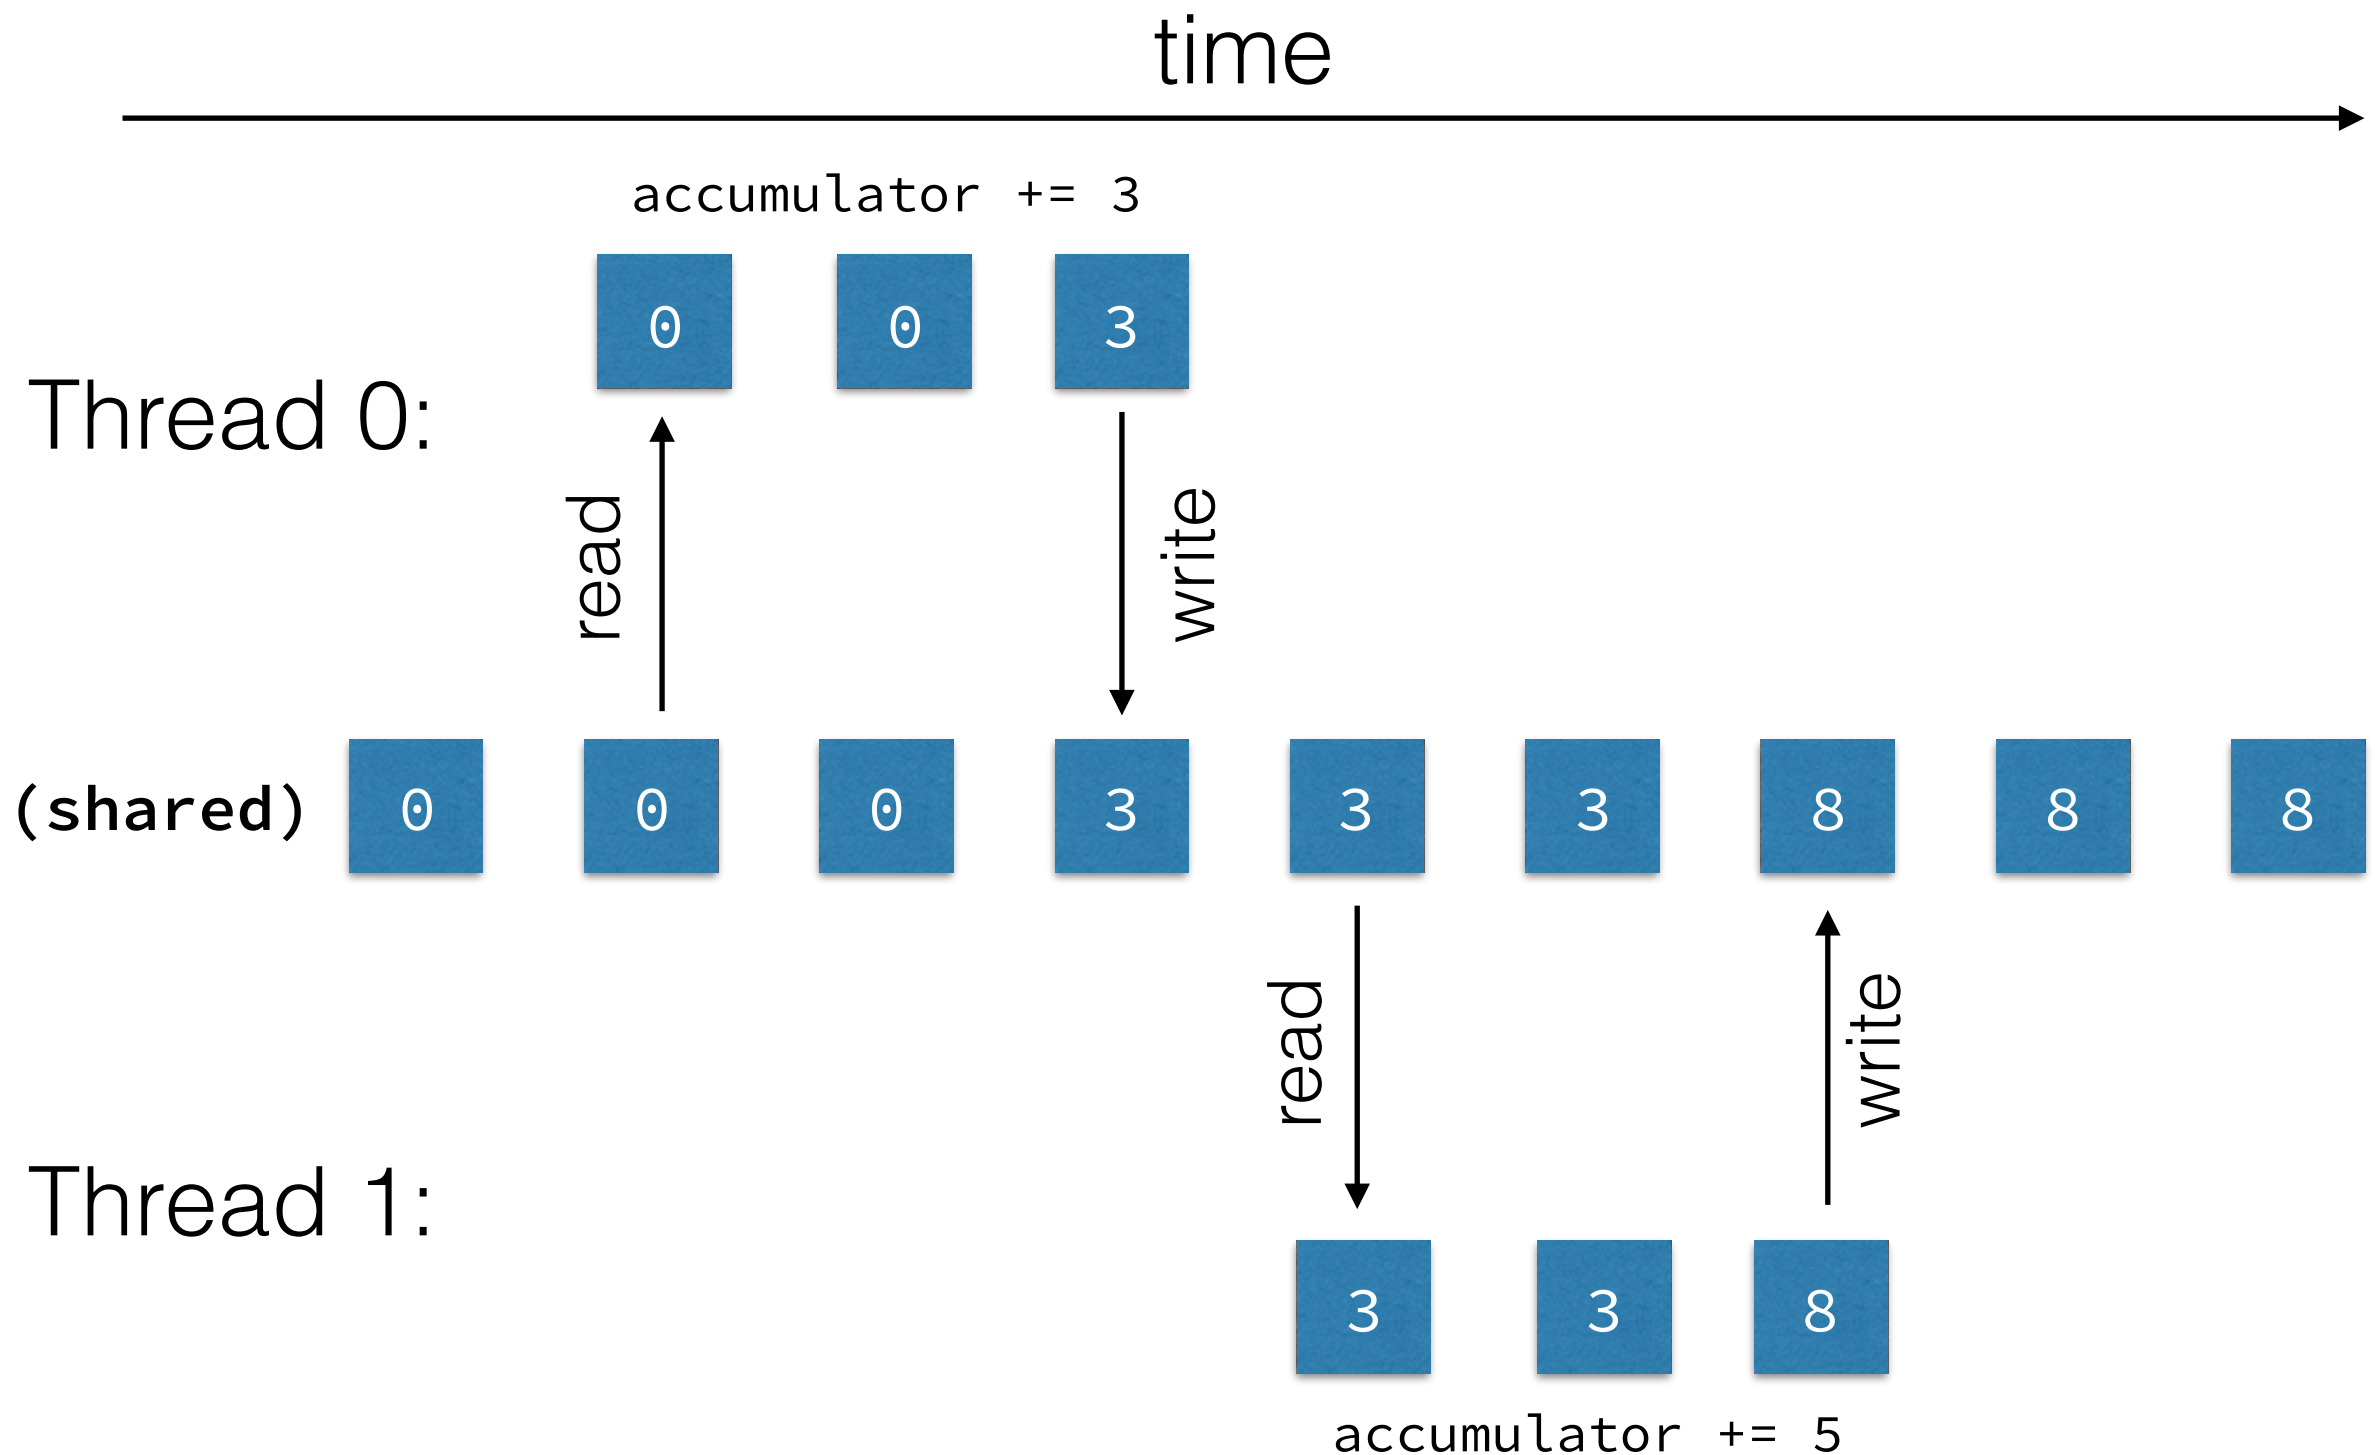
- Mutual exclusion (mutex) — only one thread at a time can enter critical region

```
double accumulator = 0;
#pragma omp parallel
{
    double output;
    int thread_id = omp_get_thread_num();
    output = big_calculation(thread_id);

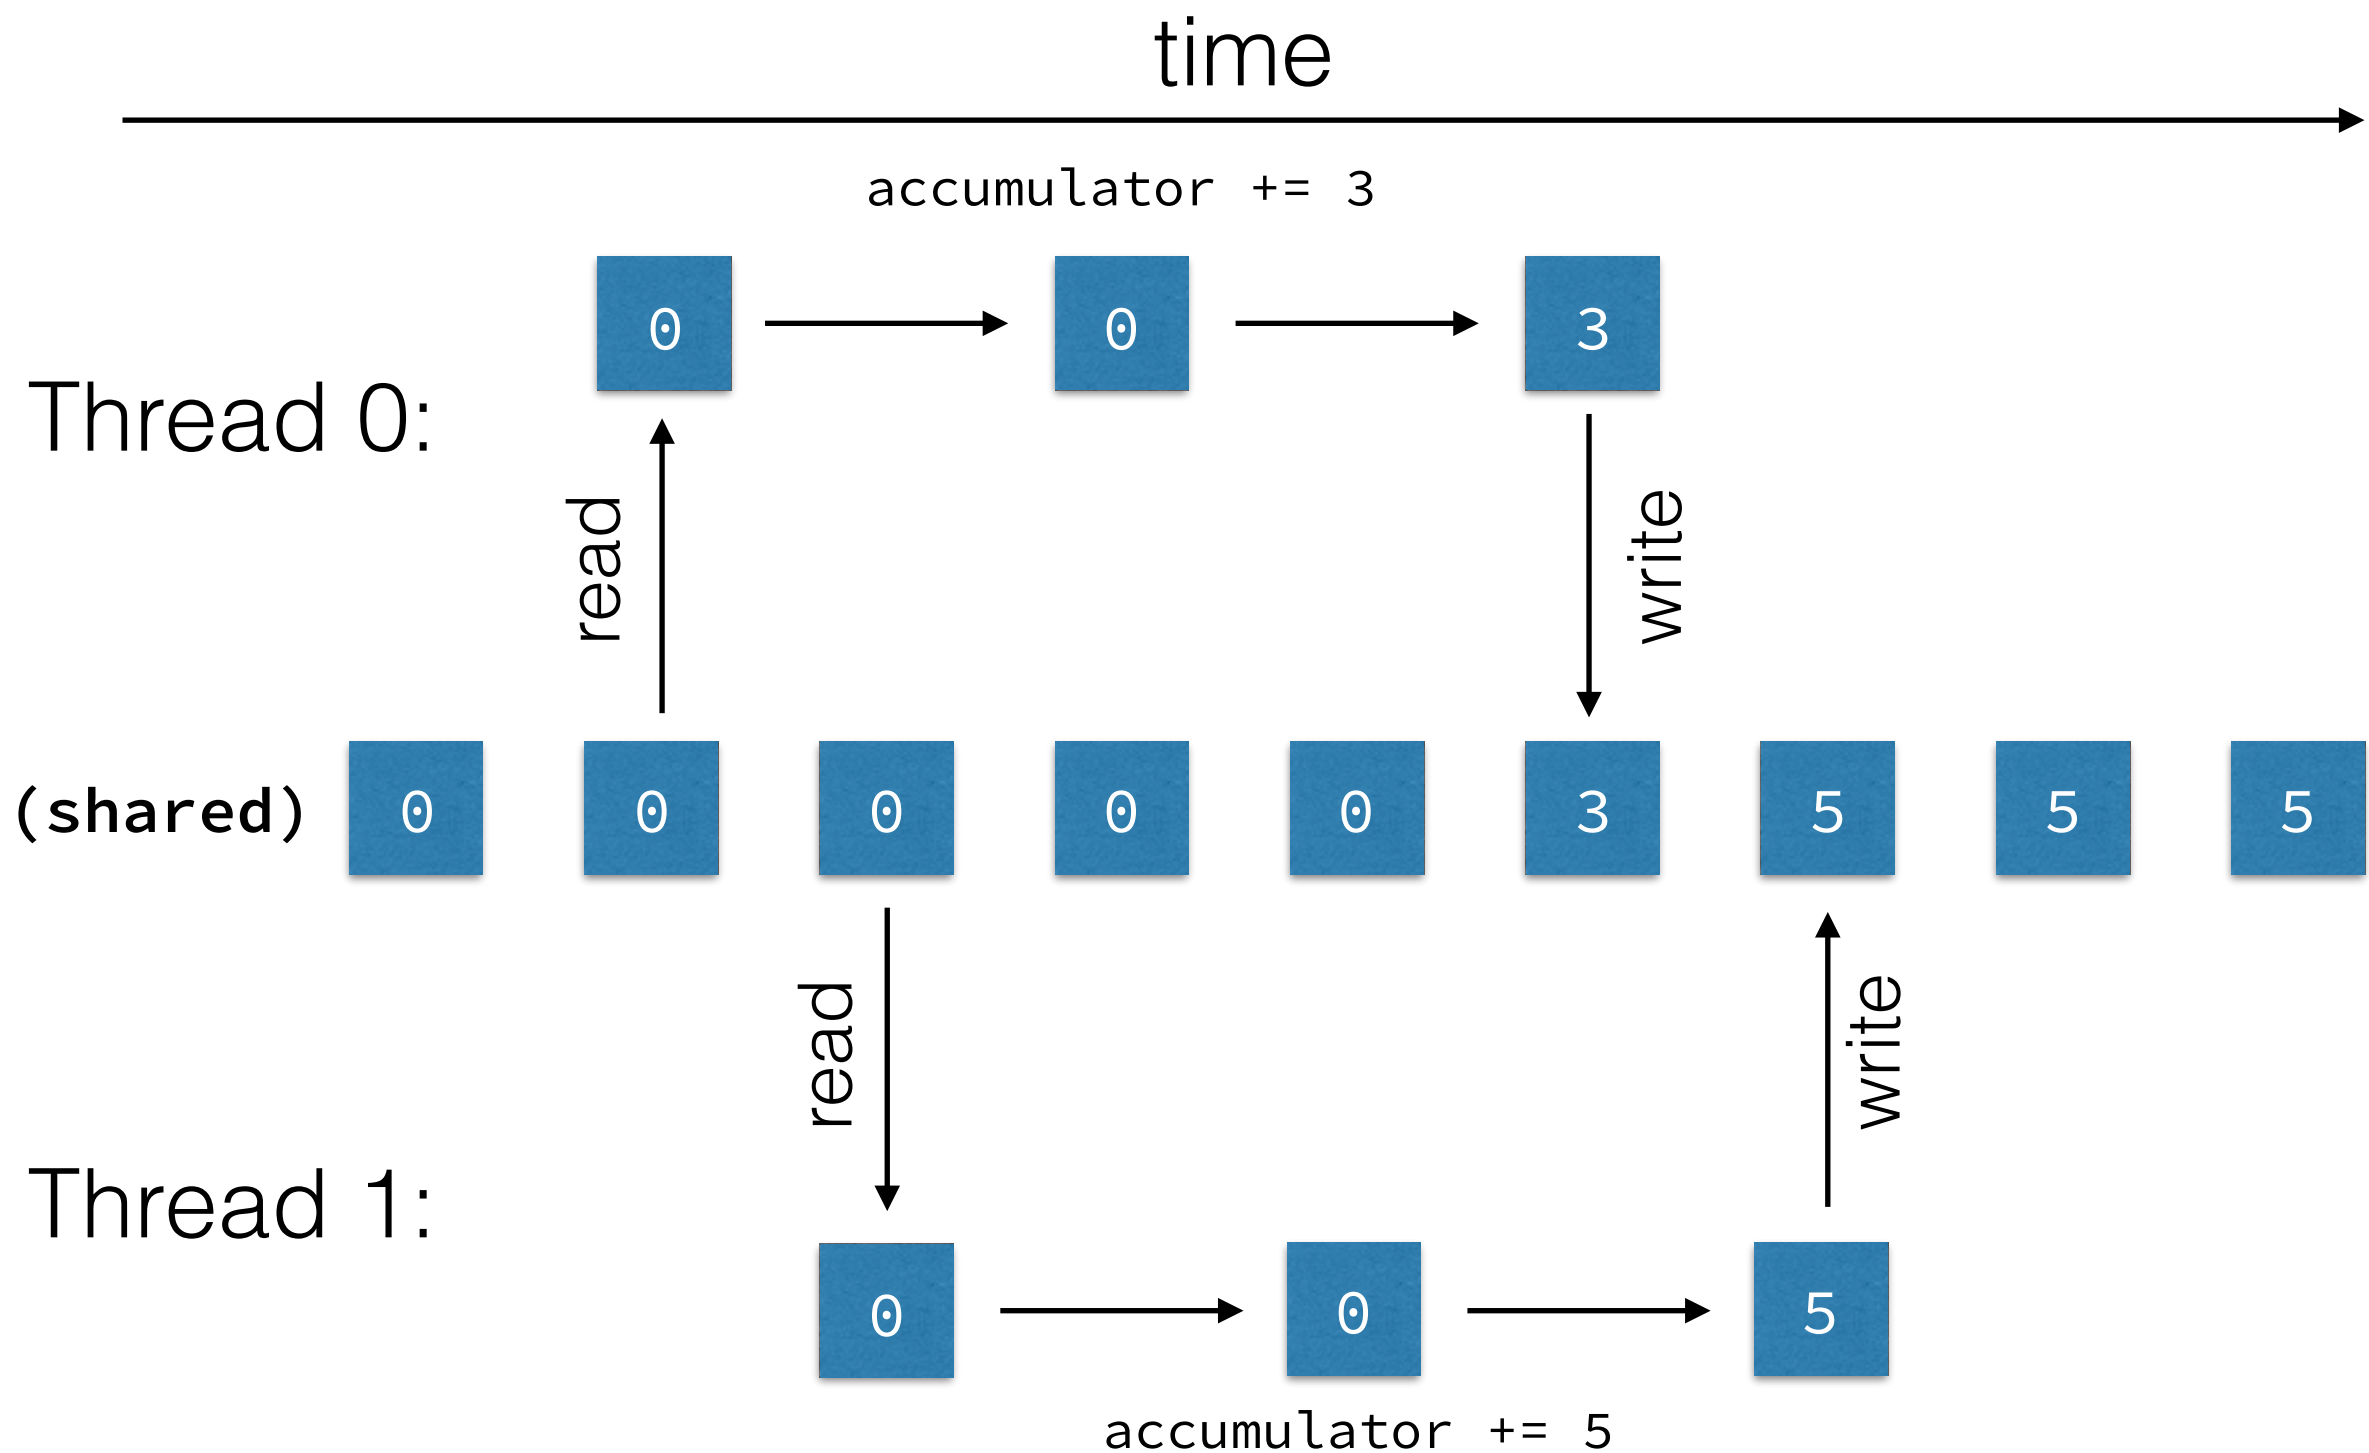
    accumulator += output;
}
```

“Race condition” — multiple threads with desynchronized read / writes

Synchronization - Critical



Synchronization - Critical



Synchronization - Critical

- Mutual exclusion (mutex) — only one thread at a time can enter critical region

```
double accumulator = 0;
#pragma omp parallel
{
    double output;
    int thread_id = omp_get_thread_num();
    output = big_calculation(thread_id);

    #pragma omp critical
    accumulator += output;
}
```

“Only one thread can execute following line at one time”

Synchronization - Atomic

- Mutex — but only for updates of memory locations

- statement inside atomic must be of form:

`shared_mem_loc BINOP= expression`

e.g. `accumulator += output;`
`accumulator *= output;`

- in-place operations also allowed:

`++accumulator;`
`accumulator--;`

Synchronization - Atomic

- “Atomic” is used in other languages for similar constructs

```
double accumulator = 0;
#pragma omp parallel
{
    double output;
    int thread_id = omp_get_thread_num();
    output = big_calculation(thread_id);

    #pragma omp atomic
    accumulator += output;
}
```

Summary of OpenMP

Concepts - Part 1

- `#pragma omp parallel` `[shared(...)]`
`[private(...)]`
`[num_threads(int)]`
- `#pragma omp for` `[schedule]`
- `#pragma omp barrier`
- `#pragma omp critical`
- `#pragma omp atomic`

Next Time

- Gradually parallelizing and improving a numerical integration calculation.