

Lecture #13 - Distributed Memory and MPI


AMath 483/583

Parallelism Grain

- **Fine Grain:** parallelize at level of individual loops, splitting work for each loop between threads
- **Coarse Grain:** split problem into large pieces and have each thread deal with one piece
 - may need to sync info at some points
 - similar to MPI


Shared Memory

- OpenMP - only shared memory environments
- single address space, multiple threads



Distributed Memory

- Each processor / machine has separate memory
 - “processes” - each have separate address space
 - process communication must be explicit



MPI - Message Passing Interface

- **Message Passing**
 - SIMD - Same Instruction Multiple Data
 - “Parent program” manages memory by placing data in processes
 - Data *explicitly* sent to/from processes.
- MPI = de facto standard for distributed computing

MPI - Message Passing Interface

- **Implementations**
 - OpenMPI (www.open-mpi.org)
 - MPICH — Argonne National Lab
 - MPIICC — Intel
- “MPI Standard” implemented by above compilers

Hello World

```
#include <mpi.h>
#include <stdio.h>

int main(int argc, char** argv)
{
    MPI_Init(NULL, NULL);

    int num_procs;
    MPI_Comm_size(MPI_COMM_WORLD, &num_procs);

    int proc_id;
    MPI_Comm_rank(MPI_COMM_WORLD, &proc_id);

    printf("Hello from proc %d of %d.\n",
           proc_id, num_procs);

    MPI_Finalize();
}
```

Hello World

```
#include <mpi.h>
#include <stdio.h>

int main(int argc, char** argv)
{
    MPI_Init(NULL, NULL);

    int num_procs;
    MPI_Comm_size(MPI_COMM_WORLD, &num_procs);

    int proc_id;
    MPI_Comm_rank(MPI_COMM_WORLD, &proc_id);

    printf("Hello from proc %d\n",
           proc_id, num_procs);

    MPI_Finalize();
}
```

`MPI_Init(int*, char**)`

- Required by every MPI program
- Must be first call.
- Talk about arguments later. NULL is fine for now.

Hello

```
#include <mpi.h>
#include <stdio.h>

int main(int argc, char** argv)
{
    MPI_Init(NULL, NULL);

    int num_procs;
    MPI_Comm_size(MPI_COMM_WORLD, &num_procs);

    int proc_id;
    MPI_Comm_rank(MPI_COMM_WORLD, &proc_id);

    printf("Hello from proc %d of %d.\n",
           proc_id, num_procs);

    MPI_Finalize();
}
```

`MPI_Comm_size(MPI_Comm, int*)`

- Returns size of given “*communicator*”
- *Communicator* = group of procs
- `MPI_COMM_WORLD` = default constructed by MPI

Hello World

```
#include <mpi.h>
#include <stdio.h>

int main(int argc, char**
{
    MPI_Init(NULL, NULL);

    int num_procs;
    MPI_Comm_size(MPI_COMM_WORLD

    int proc_id;
    MPI_Comm_rank(MPI_COMM_WORLD, &proc_id);

    printf("Hello from proc %d of %d.\n",
           proc_id, num_procs);

    MPI_Finalize();
}
```

`MPI_Comm_rank(MPI_Comm, int*)`

- Returns “*rank*” of proc running this code
- Communicator automatically assigns *rank*
- Primary method for identifying procs

Hello World

```
#include <mpi.h>
#include <stdio.h>

int main(int argc, char** argv)
{
    MPI_Init(NULL, NULL);

    int num_procs;
    MPI_Comm_size(MPI_COMM_WORLD, &num_procs);

    int proc_id;
    MPI_Comm_rank(MPI_COMM_WORLD, &proc_id);

    printf("Hello from process %d of %d\n",
           proc_id, num_procs);

    MPI_Finalize();
}
```

MPI_Finalize()

- Clean up MPI environment
- No MPI calls allowed after this

Hello World

```
#include  
#include
```

```
int main  
{  
    MPI_Init
```

Most arguments are passed by reference:

- Give MPI_Comm_size the address of num_procs
- MPI_Comm_size sets the value pointed to by num_procs to the number of processes
- num_procs is now equal to # of procs

```
int num_procs;  
MPI_Comm_size(MPI_COMM_WORLD, &num_procs);
```

```
int proc_id;  
MPI_Comm_rank(MPI_COMM_WORLD, &proc_id);
```

```
printf("Hello from proc %d of %d.\n",  
       proc_id, num_procs);
```

```
MPI_Finalize();  
}
```

Compiling Hello World

- Compile using `mpicc`
- Execute using `mpiexec`: specify number of processes

```
$ cd uwhpsc-2016/lectures/lecture16
```

```
$ mpicc hello.c
```

```
$ mpiexec -n 4 ./a.out
```

Demo

MPI Hello World

Key Observation

- Every process runs the same program `hello.c`
- Manage which processes perform which tasks using their rank / process id.
- Separate programs, separate data.

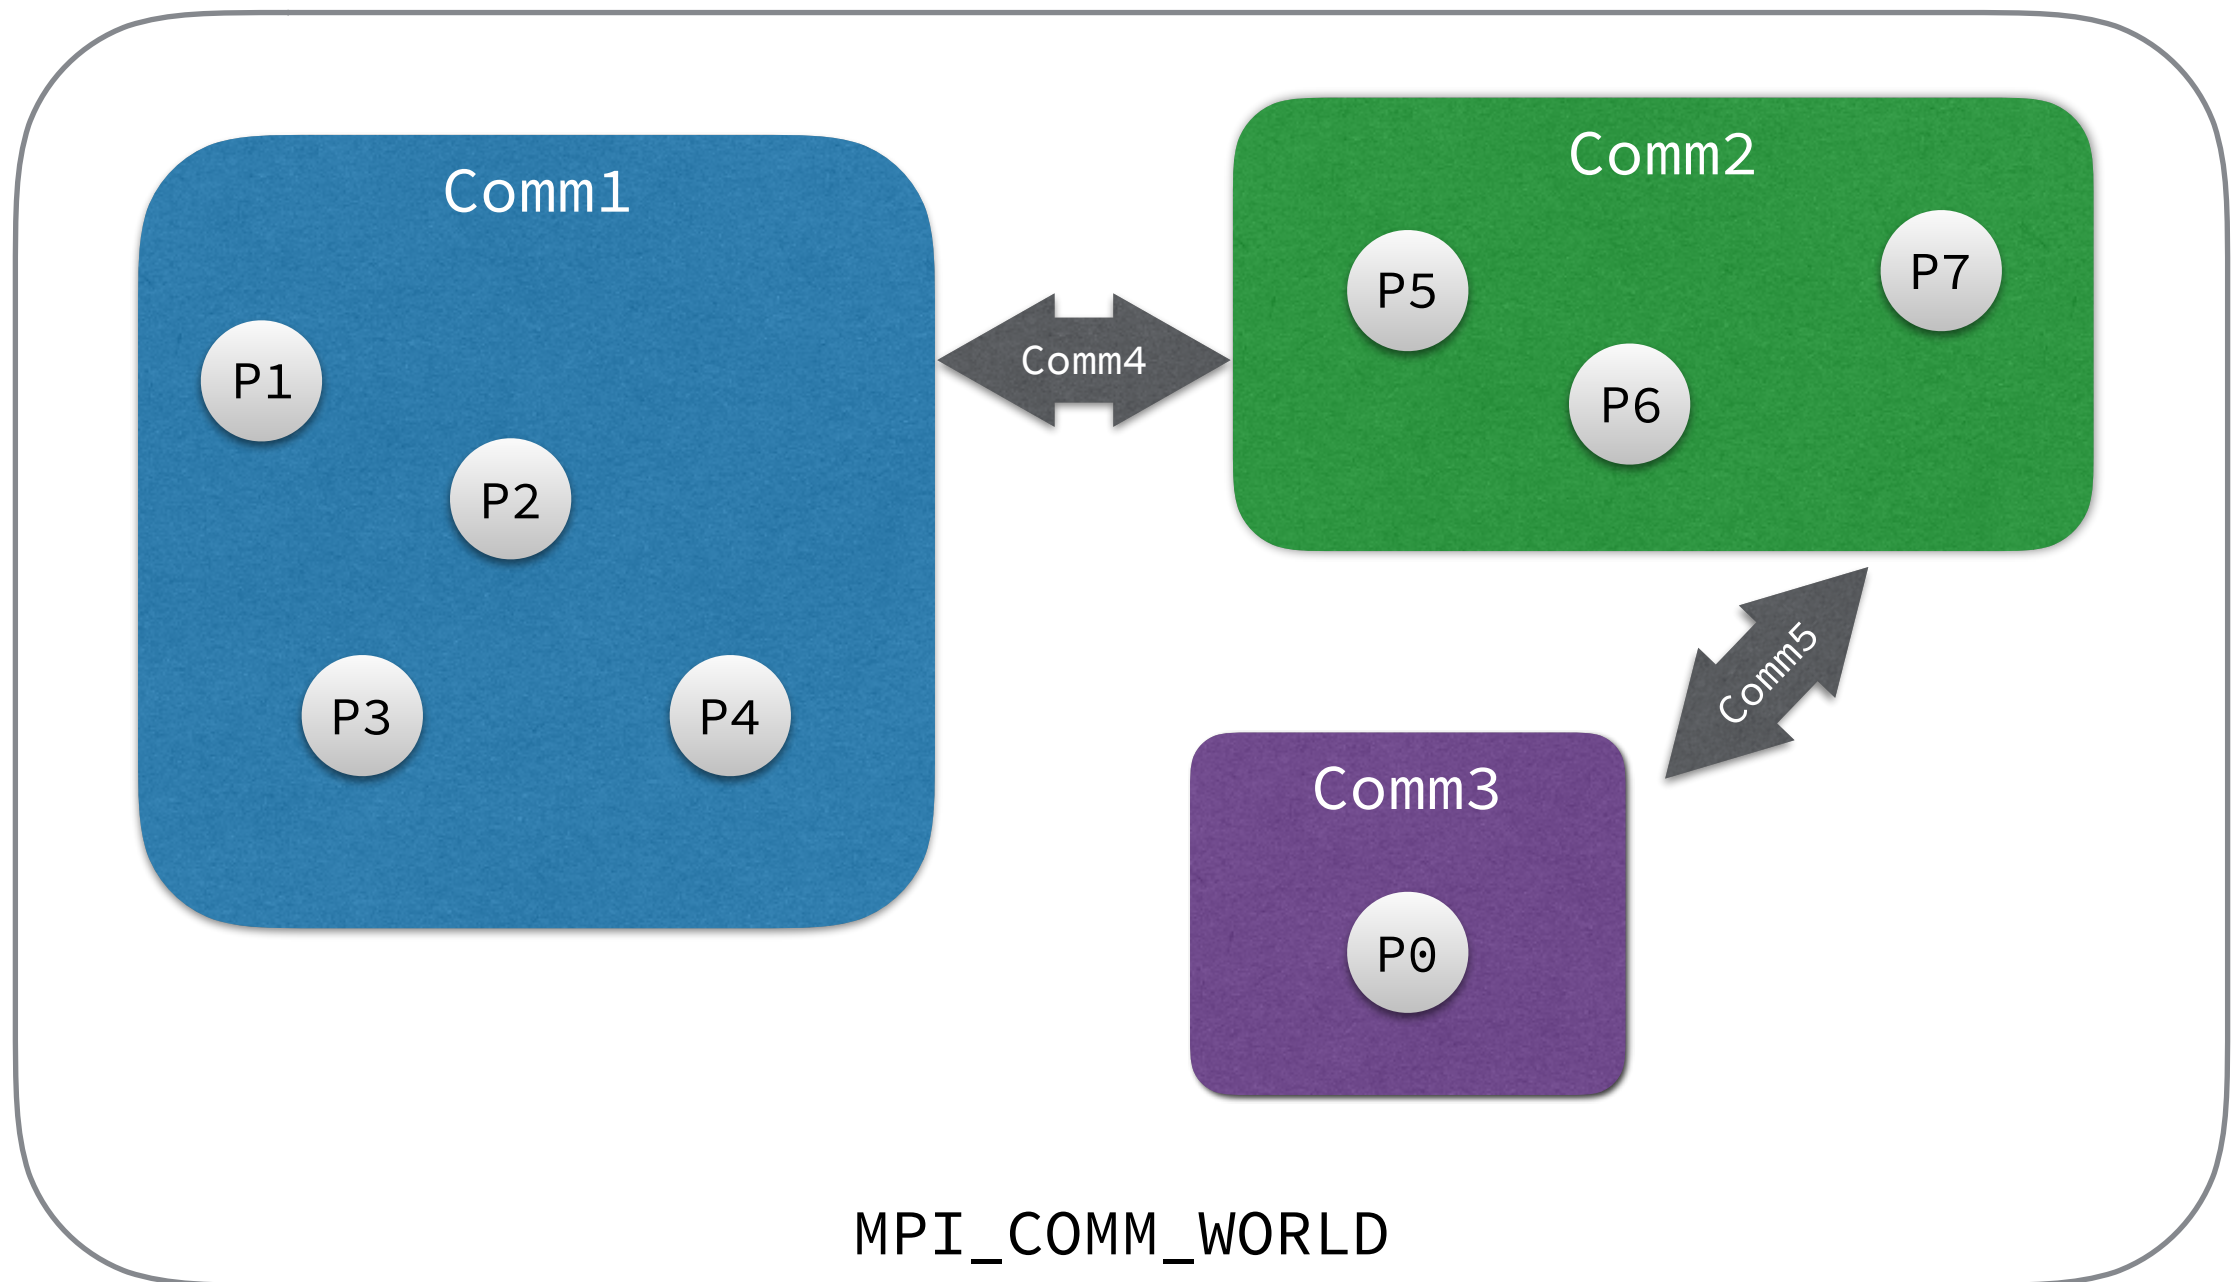
Communicators

- MPI_Comm type:
 - all communication within group of processes
 - communication takes place in some context
- MPI_COMM_WORLD — provided communicator that includes all processors by default

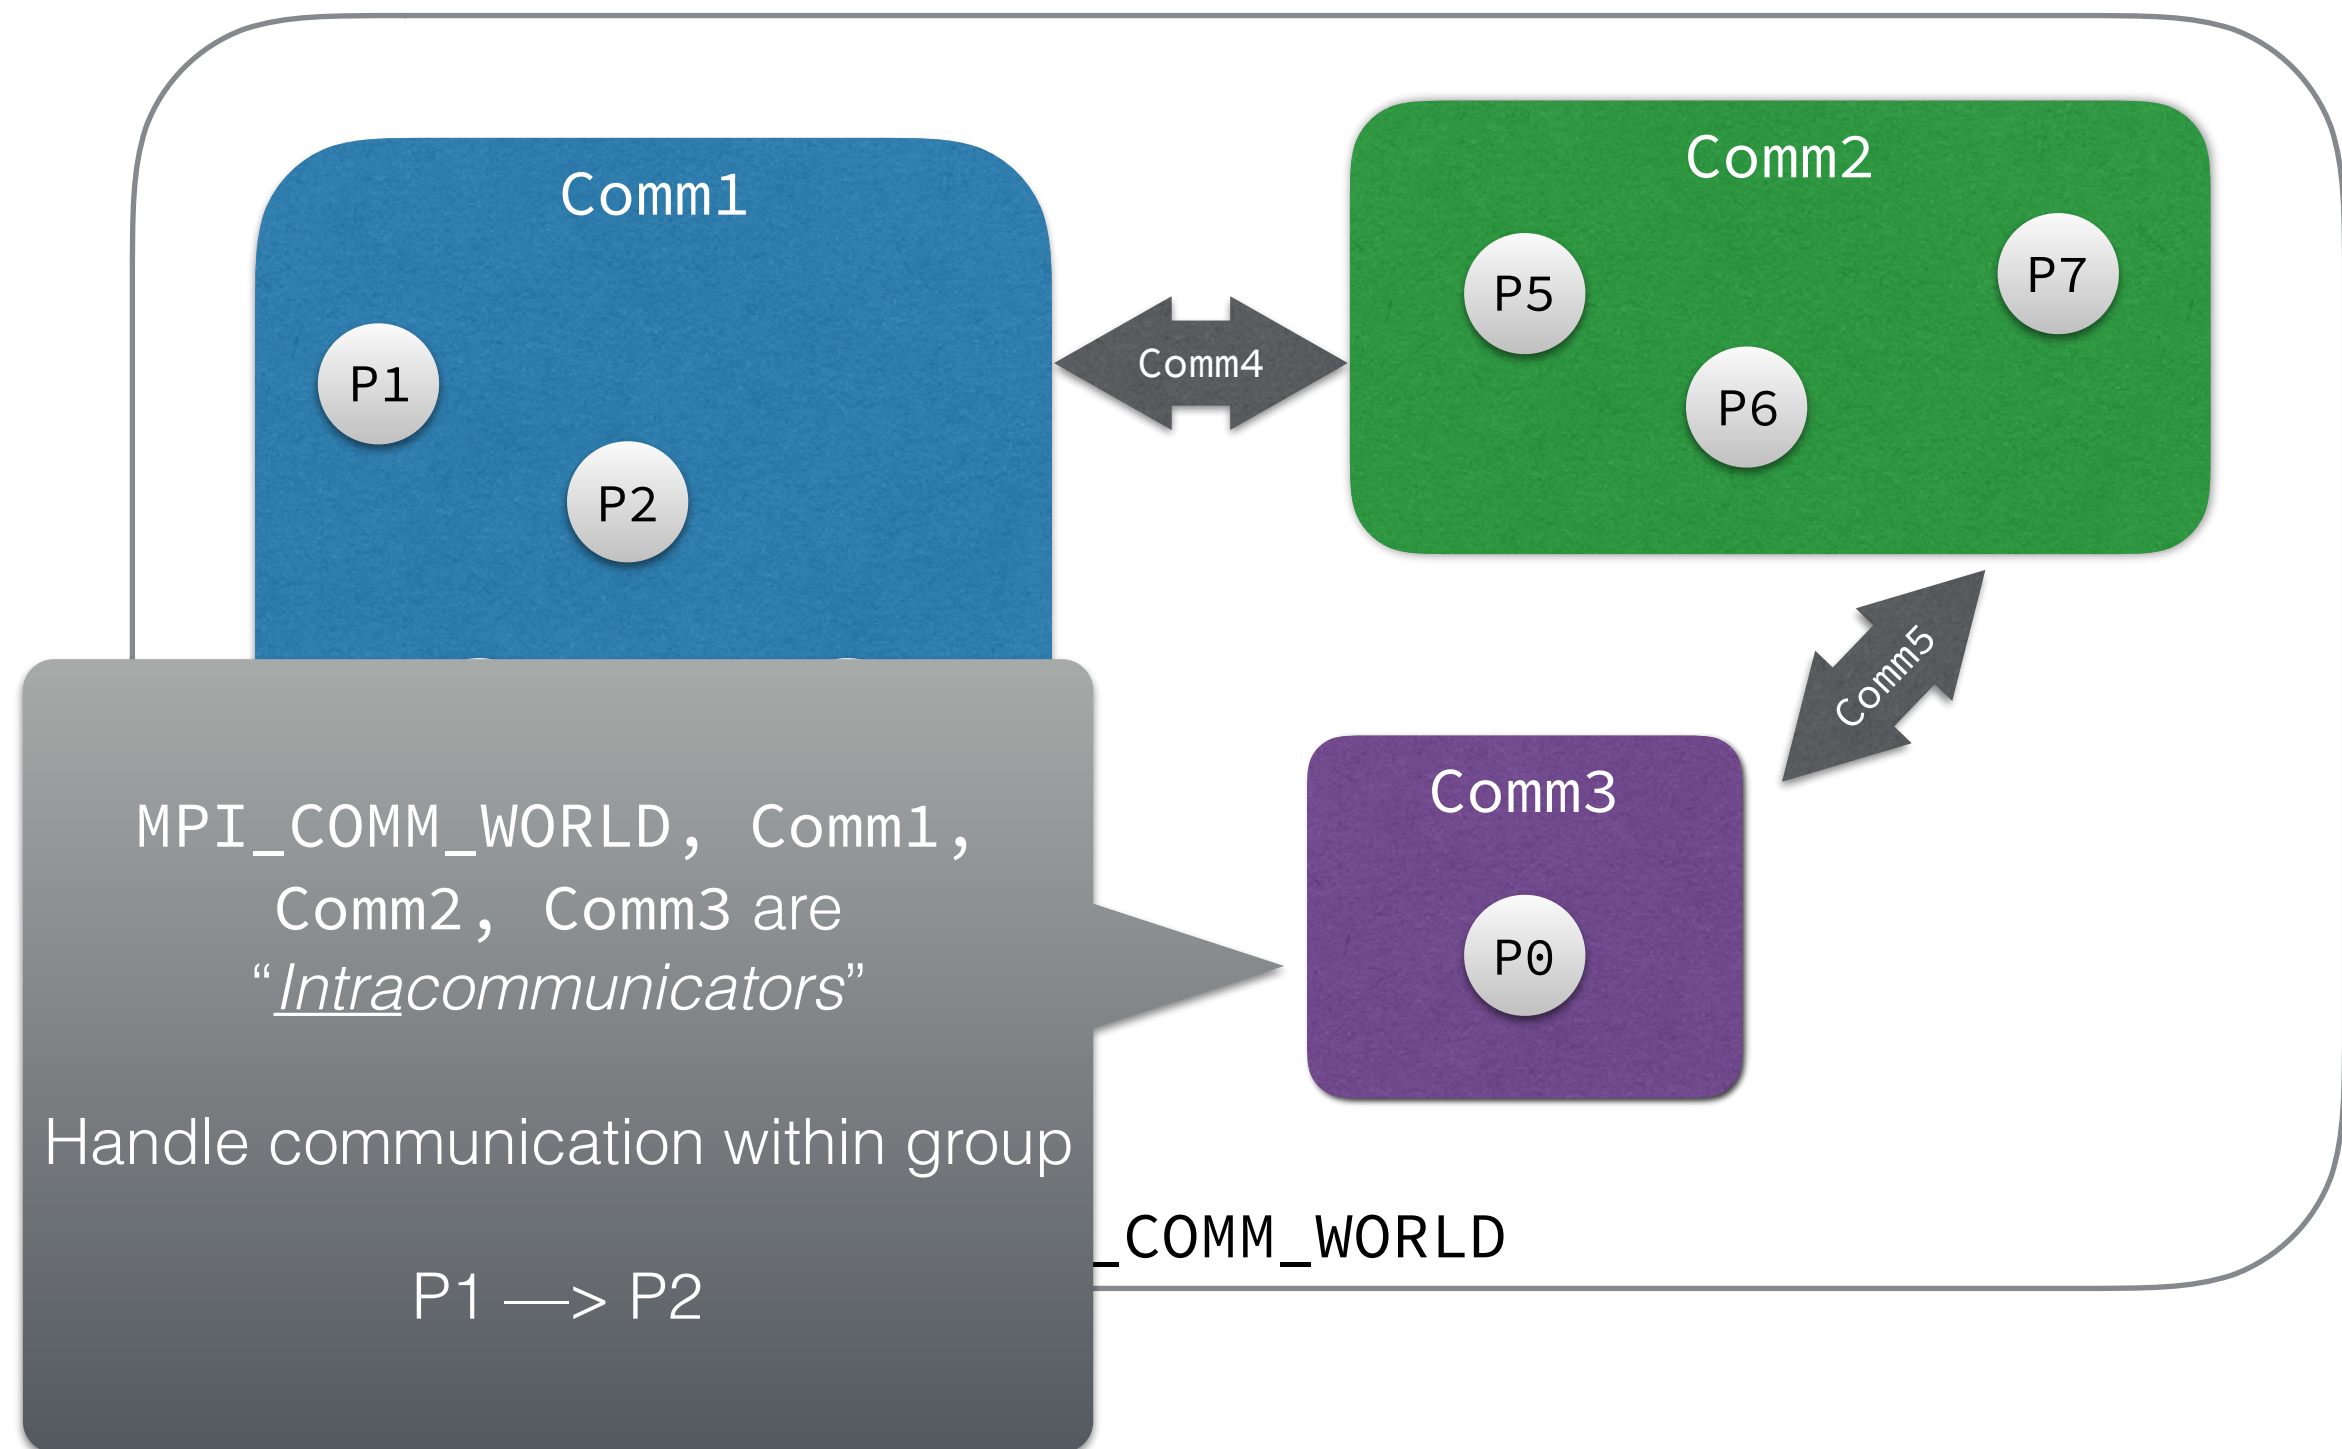
Communicators

- MPI_COMM_WORLD — provided communicator that includes all processors by default
- MPI_COMM_WORLD is sufficient for many applications
 - all communication is within a group of processes
 - communication is in some context


Communicators



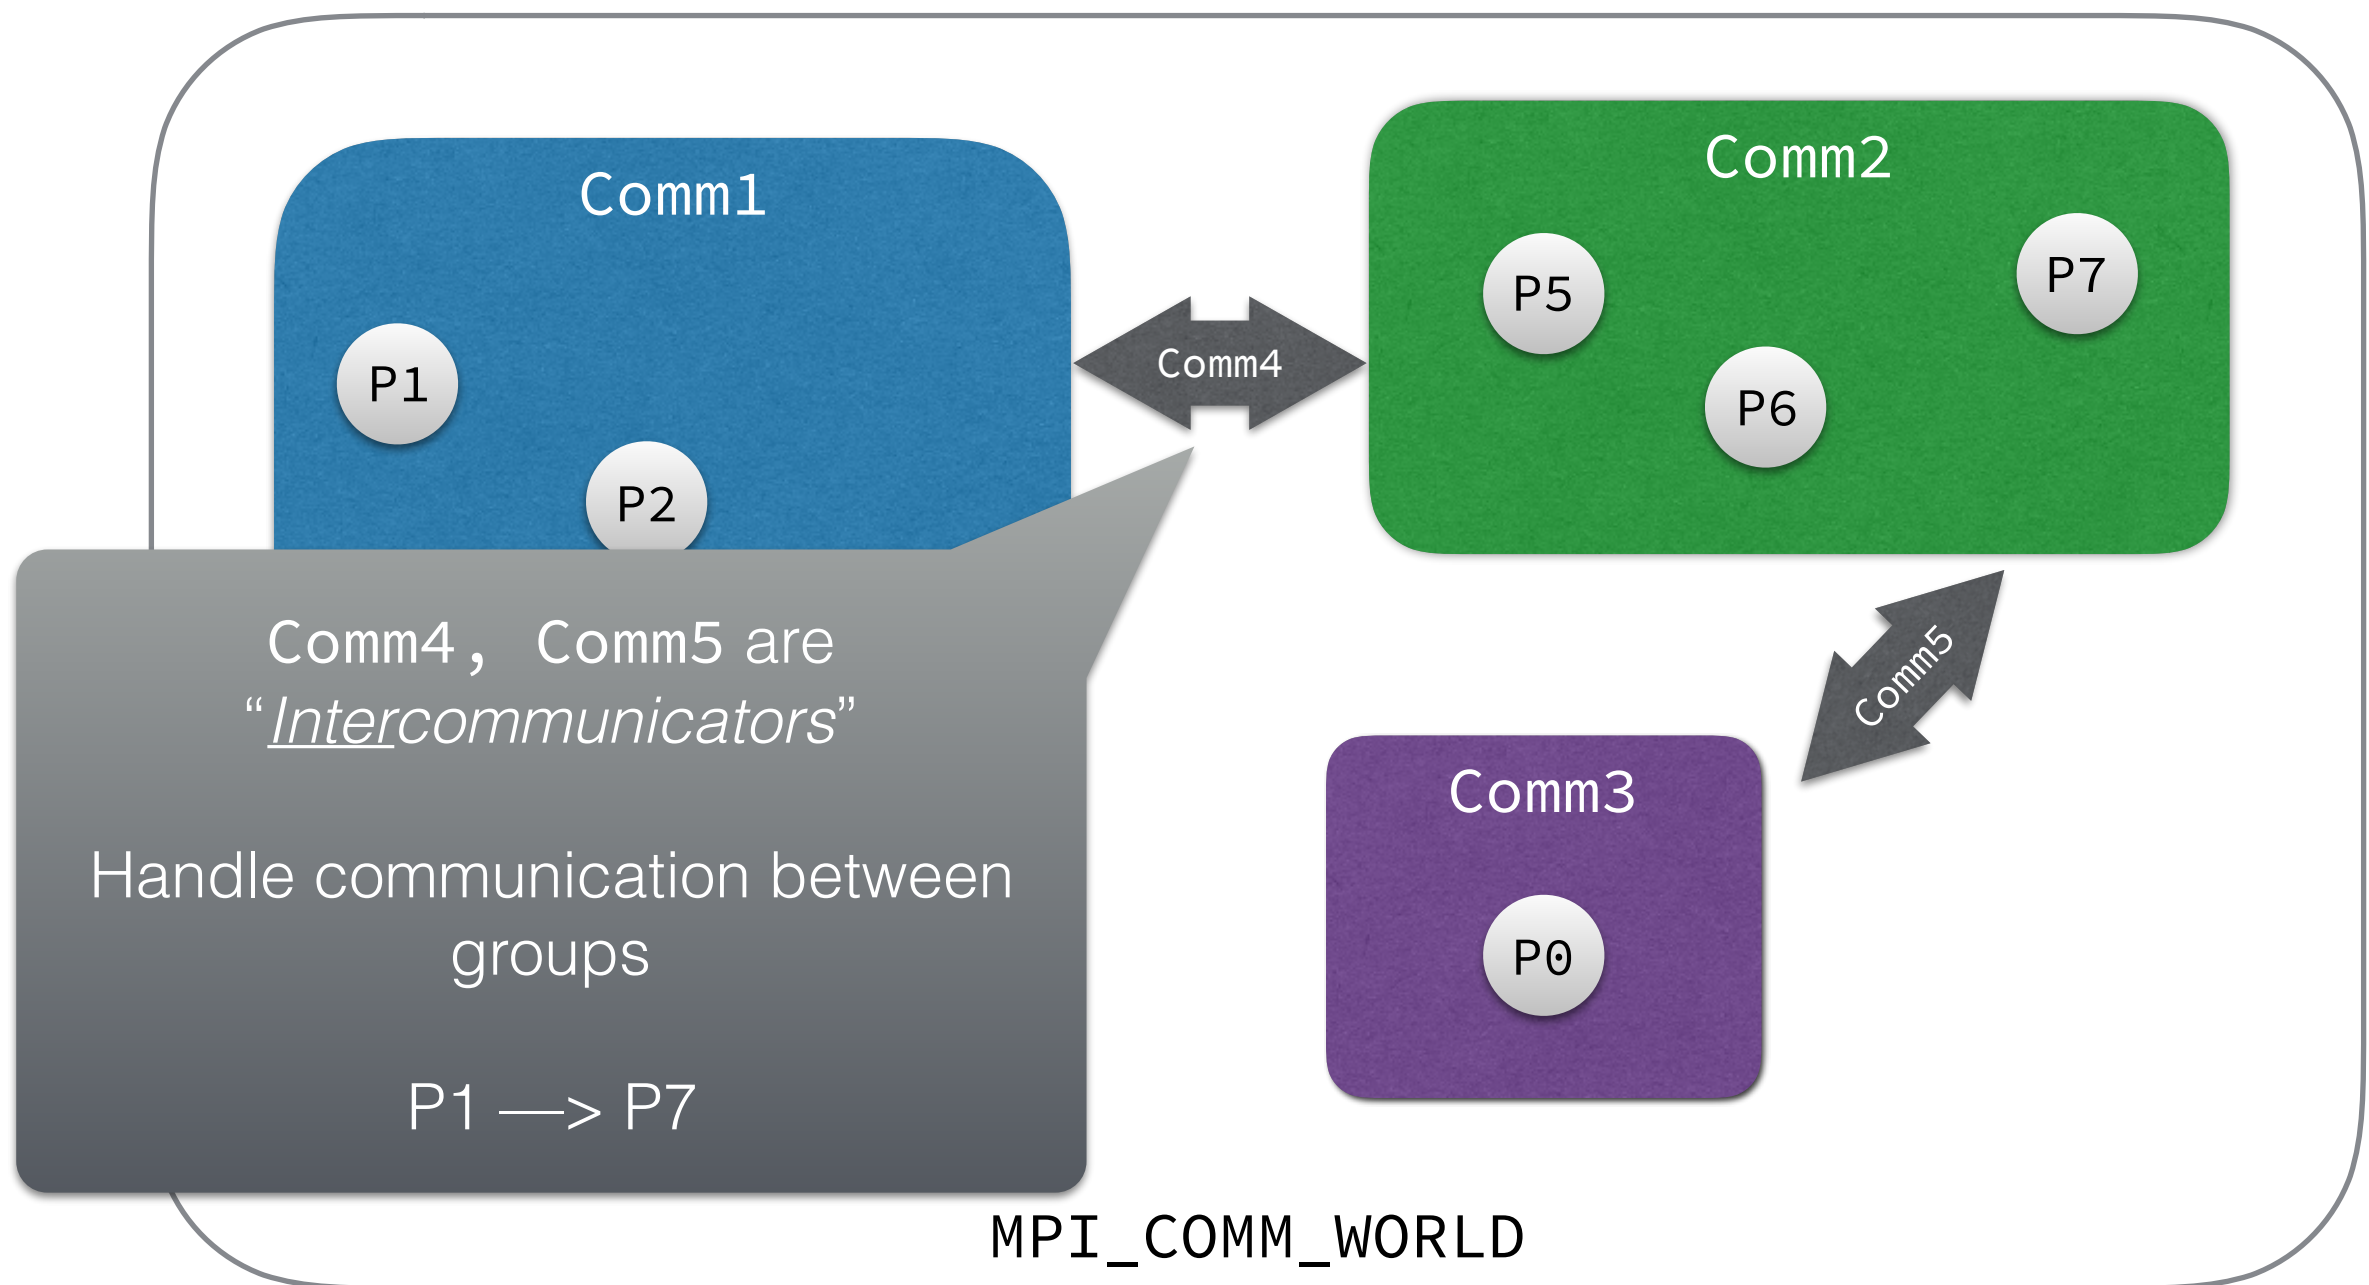
Communicators



Communicators



Communicators





MPI

- Approx. 125 MPI functions
- Many programs can be written with the following eight and MPI_COMM_WORLD
 - MPI_Init
 - MPI_Finalize
 - MPI_Comm_size
 - MPI_Comm_rank
 - MPI_Send
 - MPI_Recv
 - MPI_Bcast
 - MPI_Reduce



MPI_Send / MPI_Recv

- Basic concept behind MPI — one process sends data, another receives data



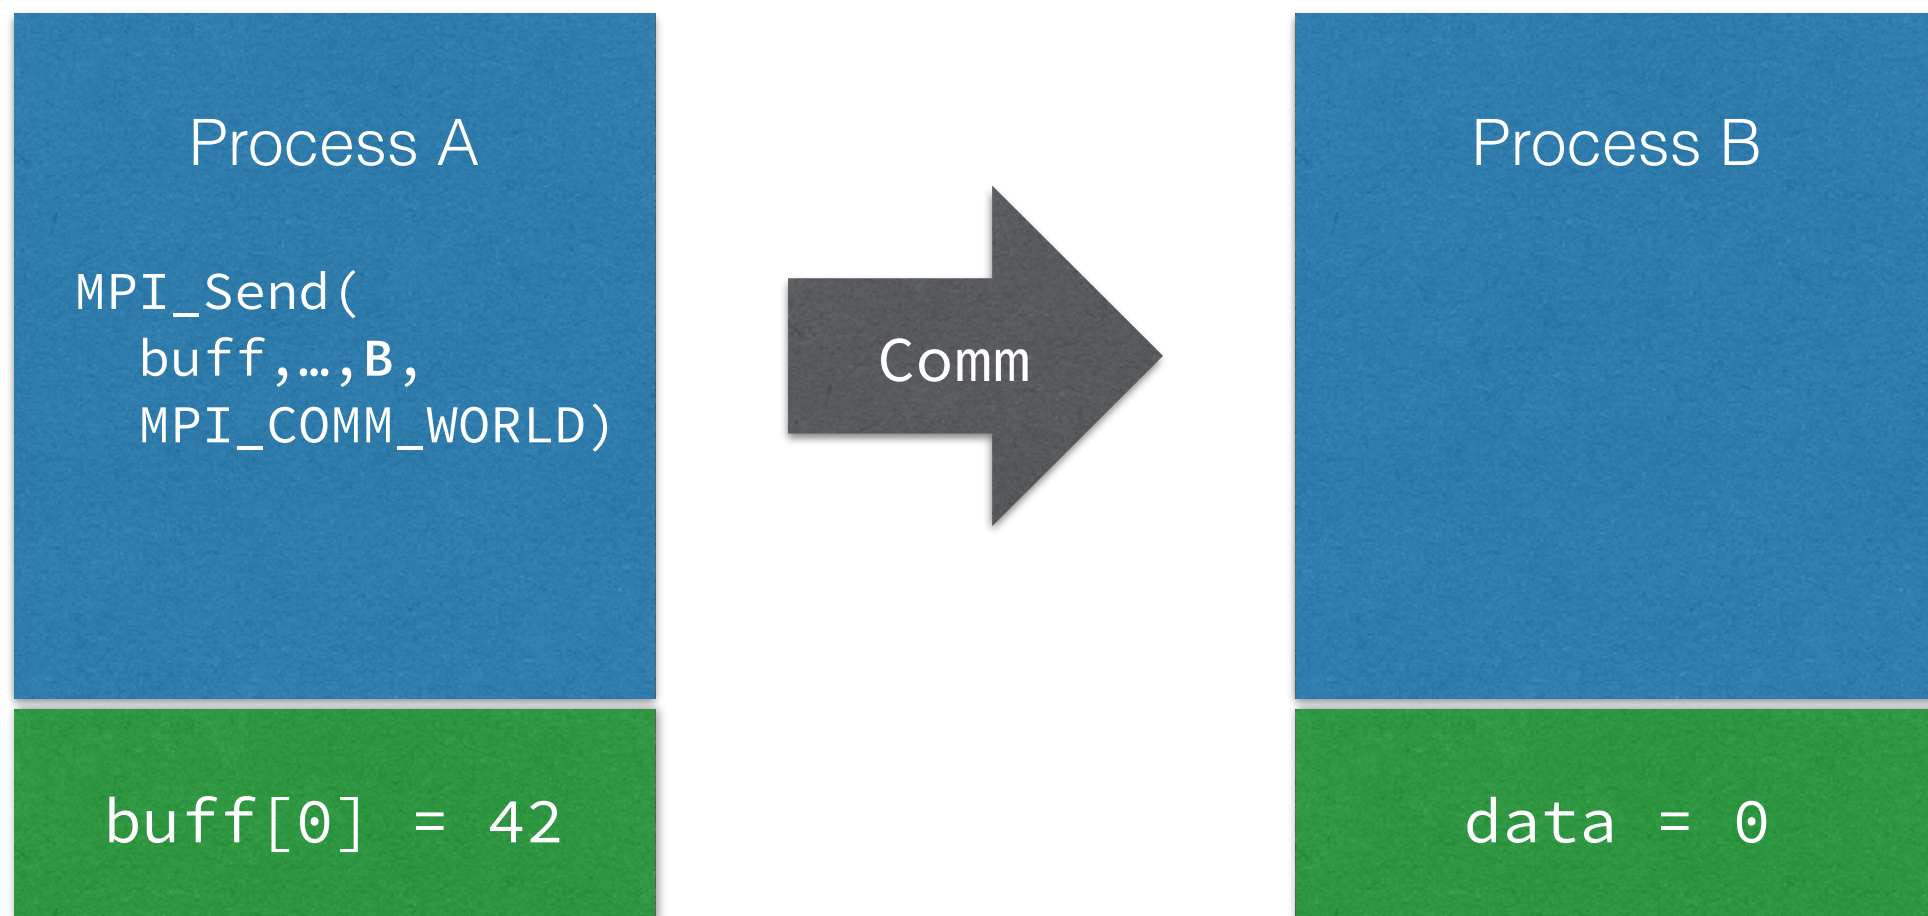
MPI_Send / MPI_Recv

- (1) Proc A packs up data to be sent into data buffer array / “*envelope*” (a pointer to data)



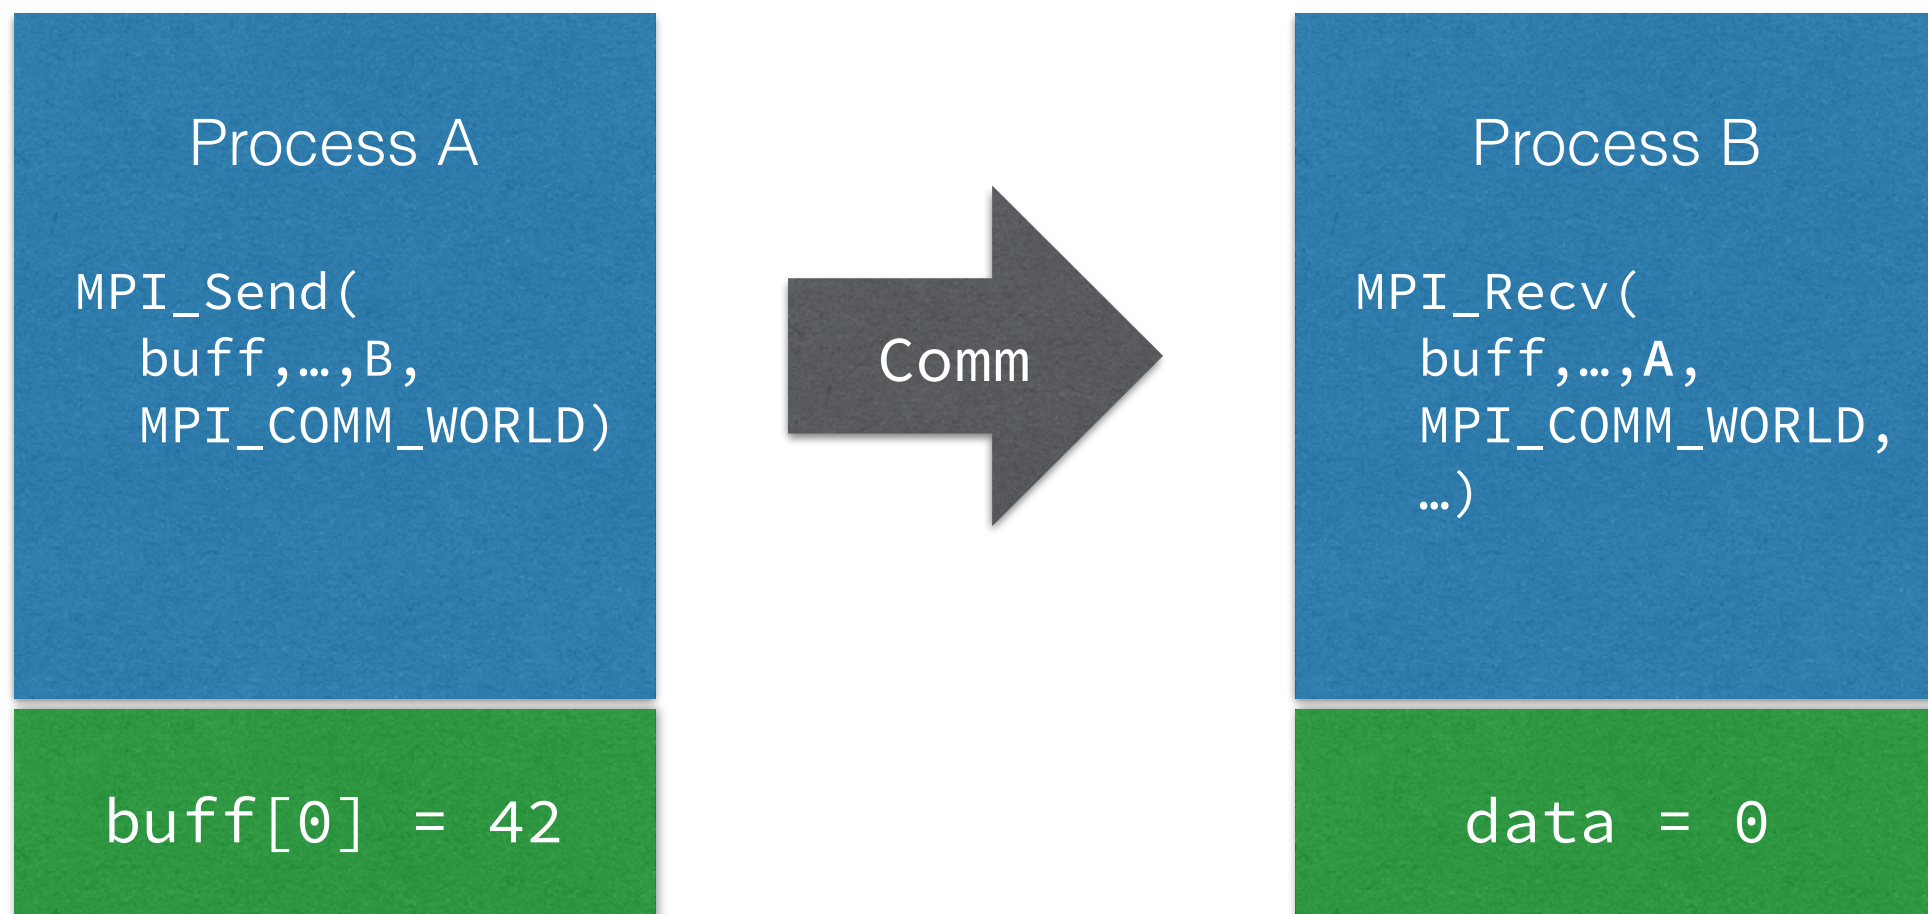
MPI_Send / MPI_Recv

- (2) Proc A announces to communicator it wants to send data buffer to Proc B with MPI_Send(...)



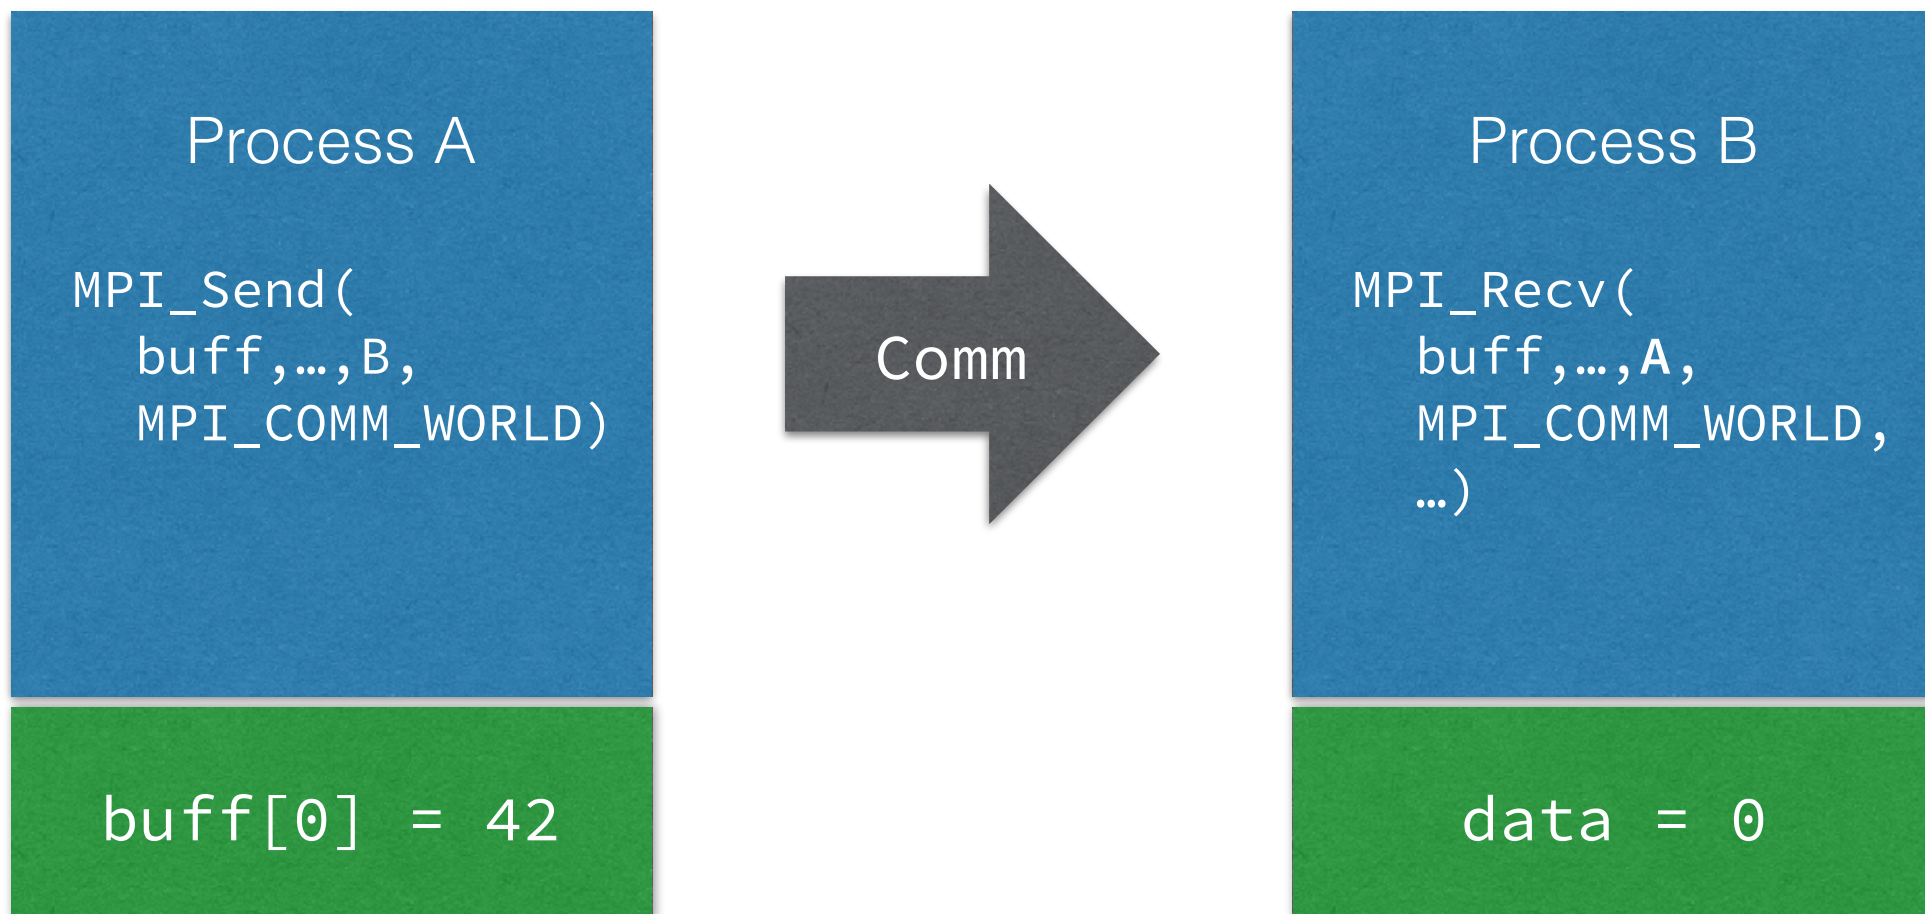
MPI_Send / MPI_Recv

- (3) Proc B acknowledges that it wants data buffer from Proc A using MPI_Recv



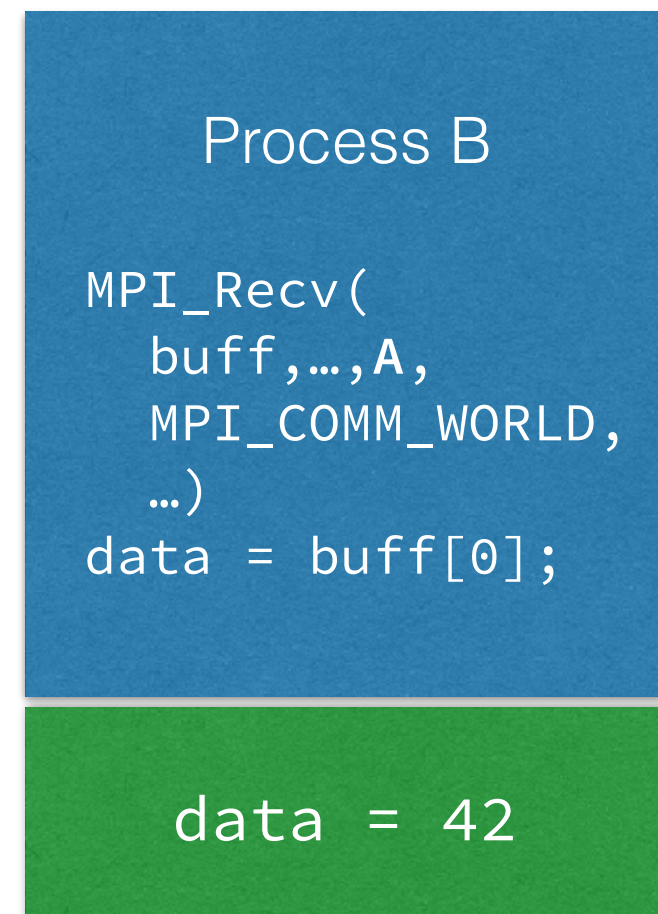

MPI_Send / MPI_Recv

- Data is not sent until Proc A reaches MPI_Send() and Proc B reaches corresponding MPI_Recv()*



MPI_Send / MPI_Recv

- (4) Proc B receives data buffer and stores / uses information



MPI_Send / MPI_Recv

```
MPI_Send(  
    void* data,  
    int count,  
    MPI_Datatype datatype,  
    int destination,  
    int tag,  
    MPI_Comm communicator)
```

```
MPI_Recv(  
    void* data,  
    int count,  
    MPI_Datatype datatype,  
    int source,  
    int tag,  
    MPI_Comm communicator,  
    MPI_Status* status)
```


MPI_Send / MPI_Recv

MPI_Send(
void* data,
int count,
MPI_Datatype datatype,
int destination,
int tag,
MPI_Comm communicator)

Pointer to some location in memory (data buffer)

Size / length of data buffer

Type of data buffer (MPI_INT, MPI_FLOAT)

Target process / where to send data

Communicator. Interprets destination.

MPI_Recv(
void* data,
int count,
MPI_Datatype datatype,
int source,
int tag,
MPI_Comm communicator,
MPI_Status* status)

MPI_Comm communicator,
MPI_Status* status)

MPI_Send / MPI_Recv

MPI_Send(

Where to store incoming data.

int count,

Where to listen for incoming data.

int destination,

int tag,

MPI_Comm communicator)

Status message (for error handling)

MPI_Recv(

void* data,

int count,

MPI_Datatype datatype,

int source,

int tag,

MPI_Comm communicator,

MPI_Status* status)

MPI_Datatype

MPI datatype	C equivalent
MPI_SHORT	short int
MPI_INT	int
MPI_LONG	long int
MPI_LONG_LONG	long long int
MPI_UNSIGNED_CHAR	unsigned char
MPI_UNSIGNED_SHORT	unsigned short int
MPI_UNSIGNED	unsigned int
MPI_UNSIGNED_LONG	unsigned long int
MPI_UNSIGNED_LONG_LONG	unsigned long long int
MPI_FLOAT	float
MPI_DOUBLE	double
MPI_LONG_DOUBLE	long double
MPI_BYTE	char

Demo

- `send-receive.c`
- Proc 0 will send a data packet to Proc 1
- Proc 1 reports that it has the data.

Demo

`send-receive.c`

Example: Distributed Minimum

- Find minimum across large array
- **Strategy:**
 - Proc 0 gives chunk of array to each process
 - each process finds local minimum
 - process reports result back to Proc 0

Pointer Reminder


```
int a[8];
```

42	10	-24	5	0	2135	90	4
0x00	0x04	0x08	0x0c	0x10	0x14	0x18	0x1c

```
    a == 0x00;  
    a[0] == 42;  
    &a[0] == 0x00;  
    a[2] == -24;  
    &a[2] == 0x08;
```

Pointer Reminder

```
int a[8];
```



```
MPI_Send(&a[4], 2, ...)
```

Demo

`distributed-min.c`

Example: Matmul

- $AB = C$ (all $N \times N$)

```
int i,j,k;
for (i=0; i<N; ++i) {
    for (j=0; j<N; ++j) {
        C[i*N+j] = 0;
        for (k=0; k<N; ++k)
            C[i*N+j] = A[i*N+k] * B[k*N+j];
    }
}
```

