

Lecture #15 - MPI

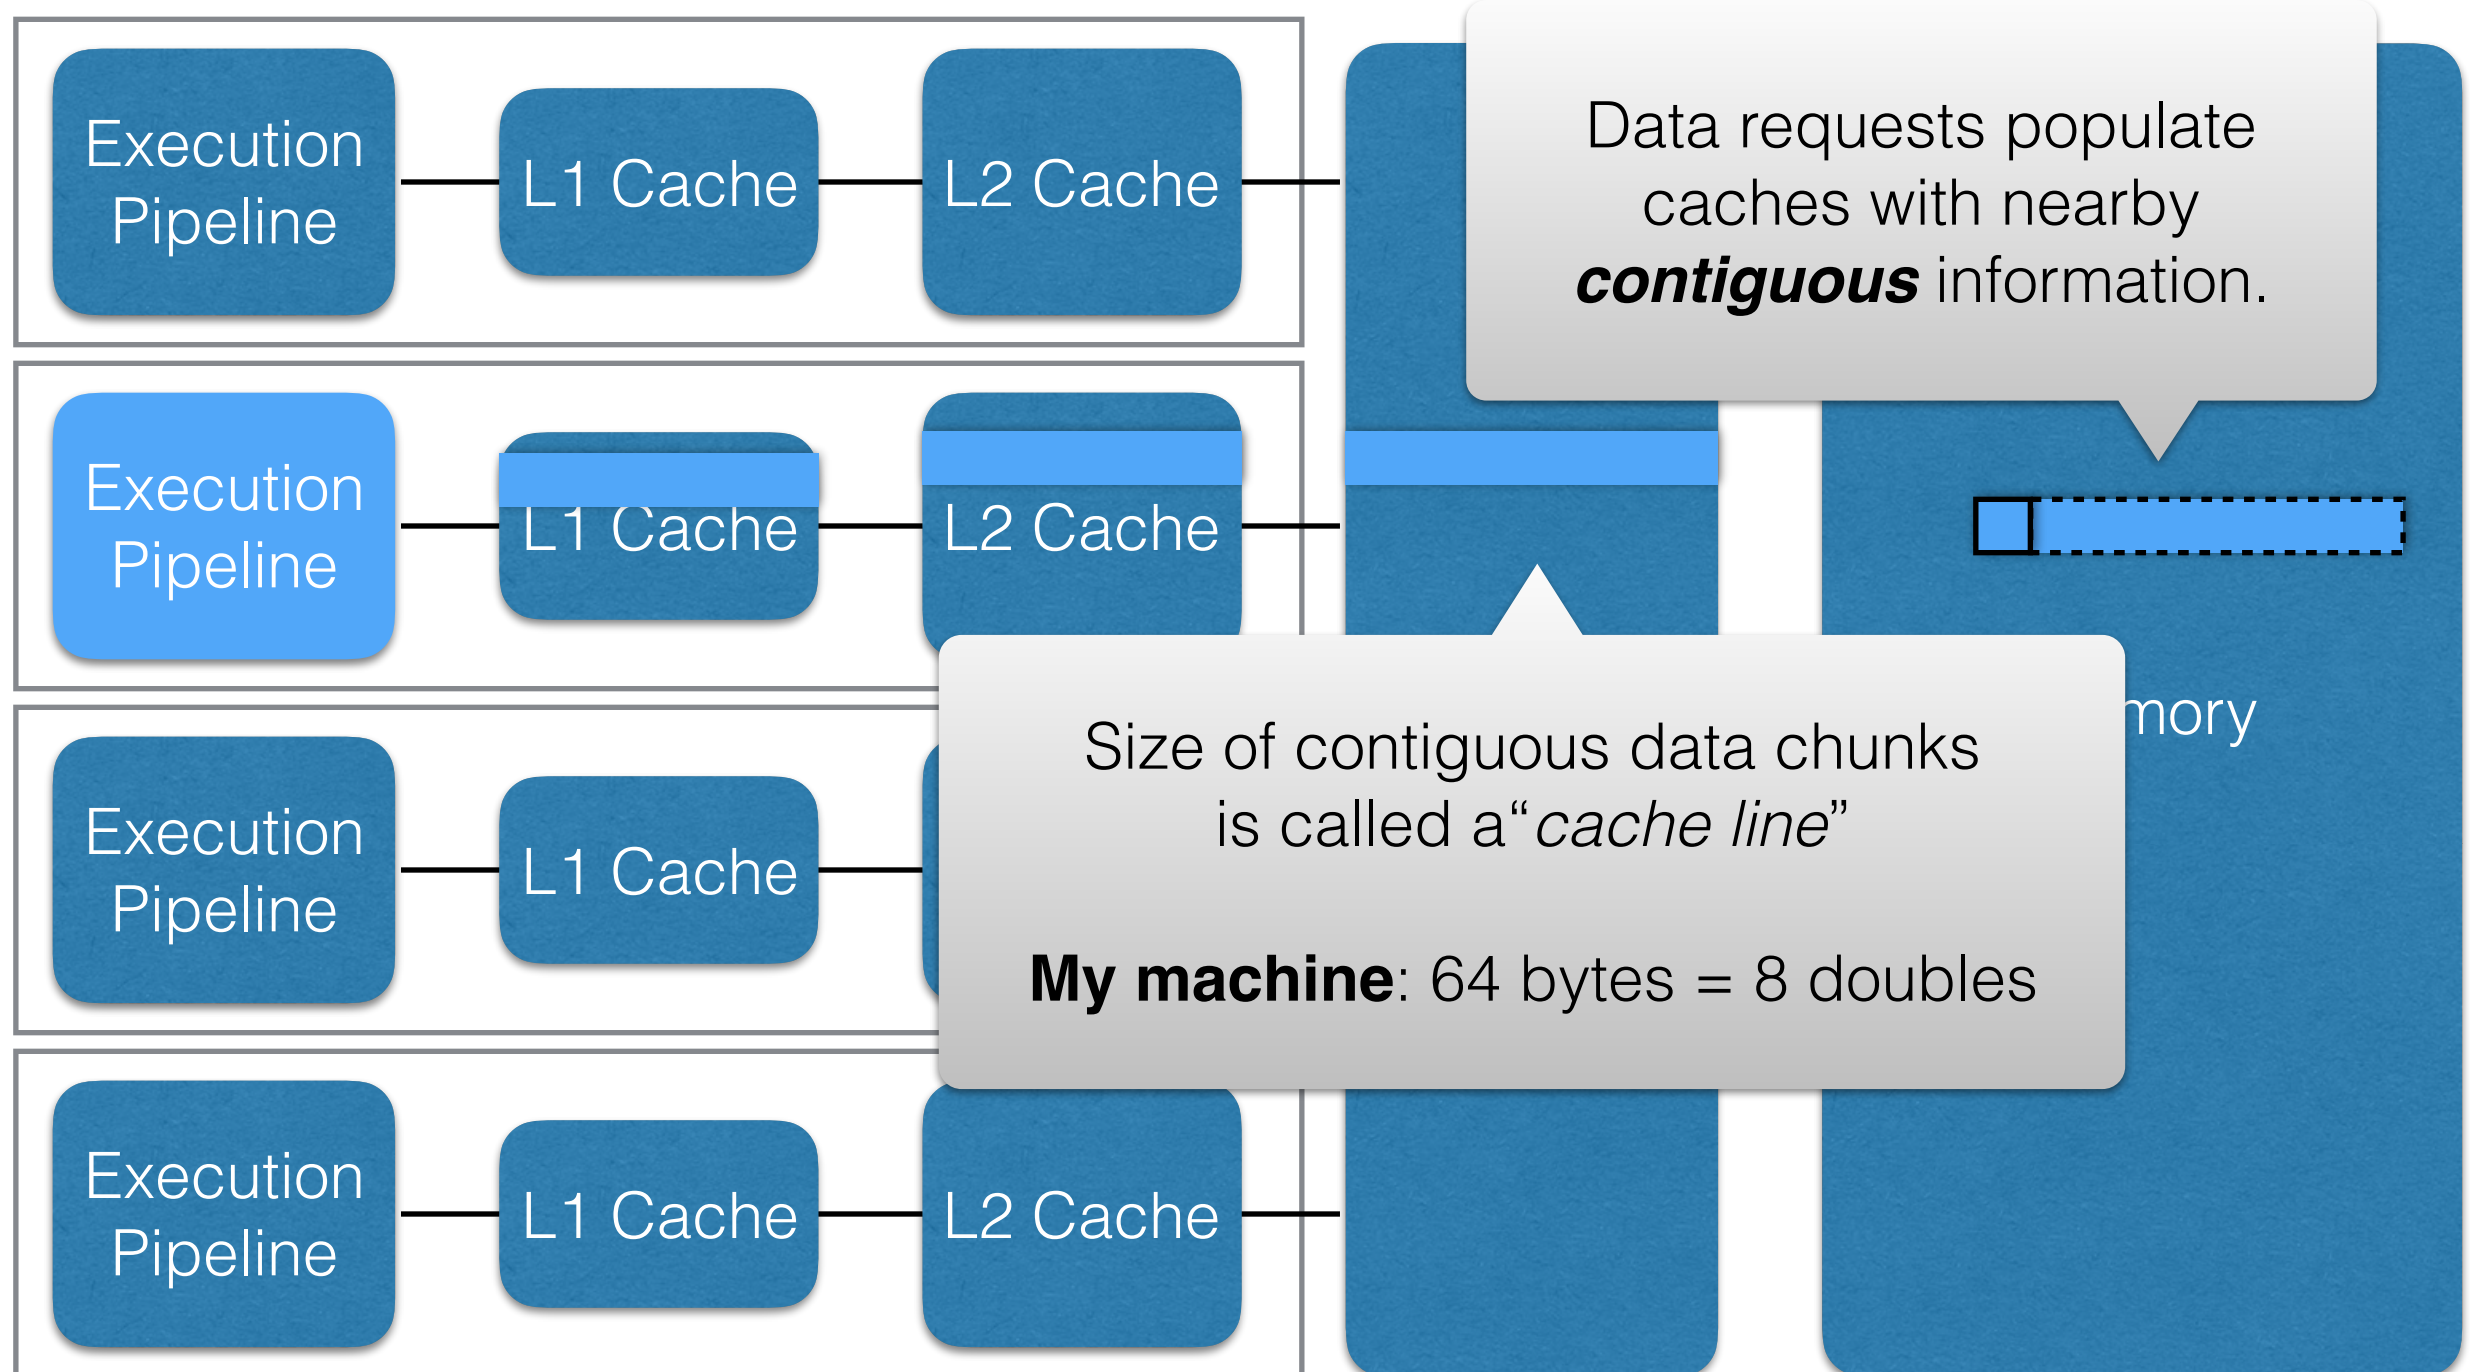
Domain Decomposition

AMath 483/583

Announcements


- Homework #3 — Due Friday (*get started if you haven't already! SMC loads are high before due date*)
- Today:
 - brief aside on cache lines
 - MPI (parallel) design patterns

Memory Layout




Cache Lines

- Imagine RAM laid out in rows of cache lines, each 8 doubles long.
- When you access a single piece of data the entire line is sent to the cache.



Cache Lines

- Why contiguity of access is good.
- Sub-optimal access patterns don't take advantage of cache lines.




Cache Lines

```
for (size_t i=0; i<N; ++i)  
    arr[i] = // input
```

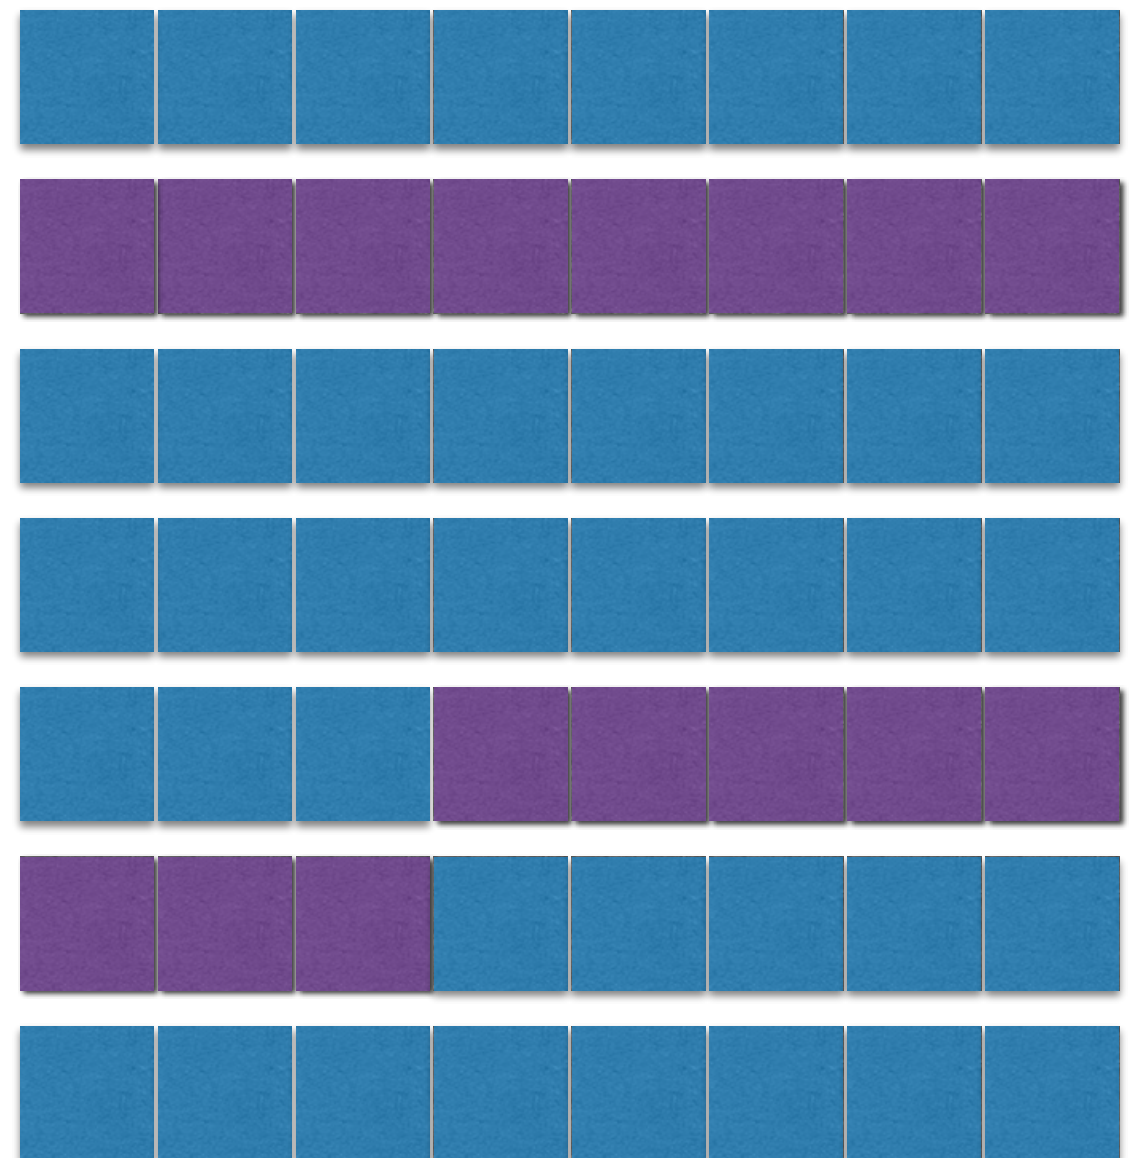
```
for (size_t i=0; i<N; i +=2 )  
    arr[i] = // input
```

```
for (size_t i=0; i<N; ++i)  
    arr[i*N+i] = // input
```



Cache Lines

- “Cache-alignment” — make most use of cache
- `posix_memalign()`
- `aligned_alloc()`
- C struct packing
- “Power of two” rule




MPI and Domain Decomposition

(Back to regular lecture)


Reminder

- MPI is used in distributed memory environments
- If working in shared memory environment, probably better off using OpenMP



Domain Decomposition

- Common “design strategy” for MPI programs
- Each process only computes over subdomain
- May need to communicate on “boundaries”



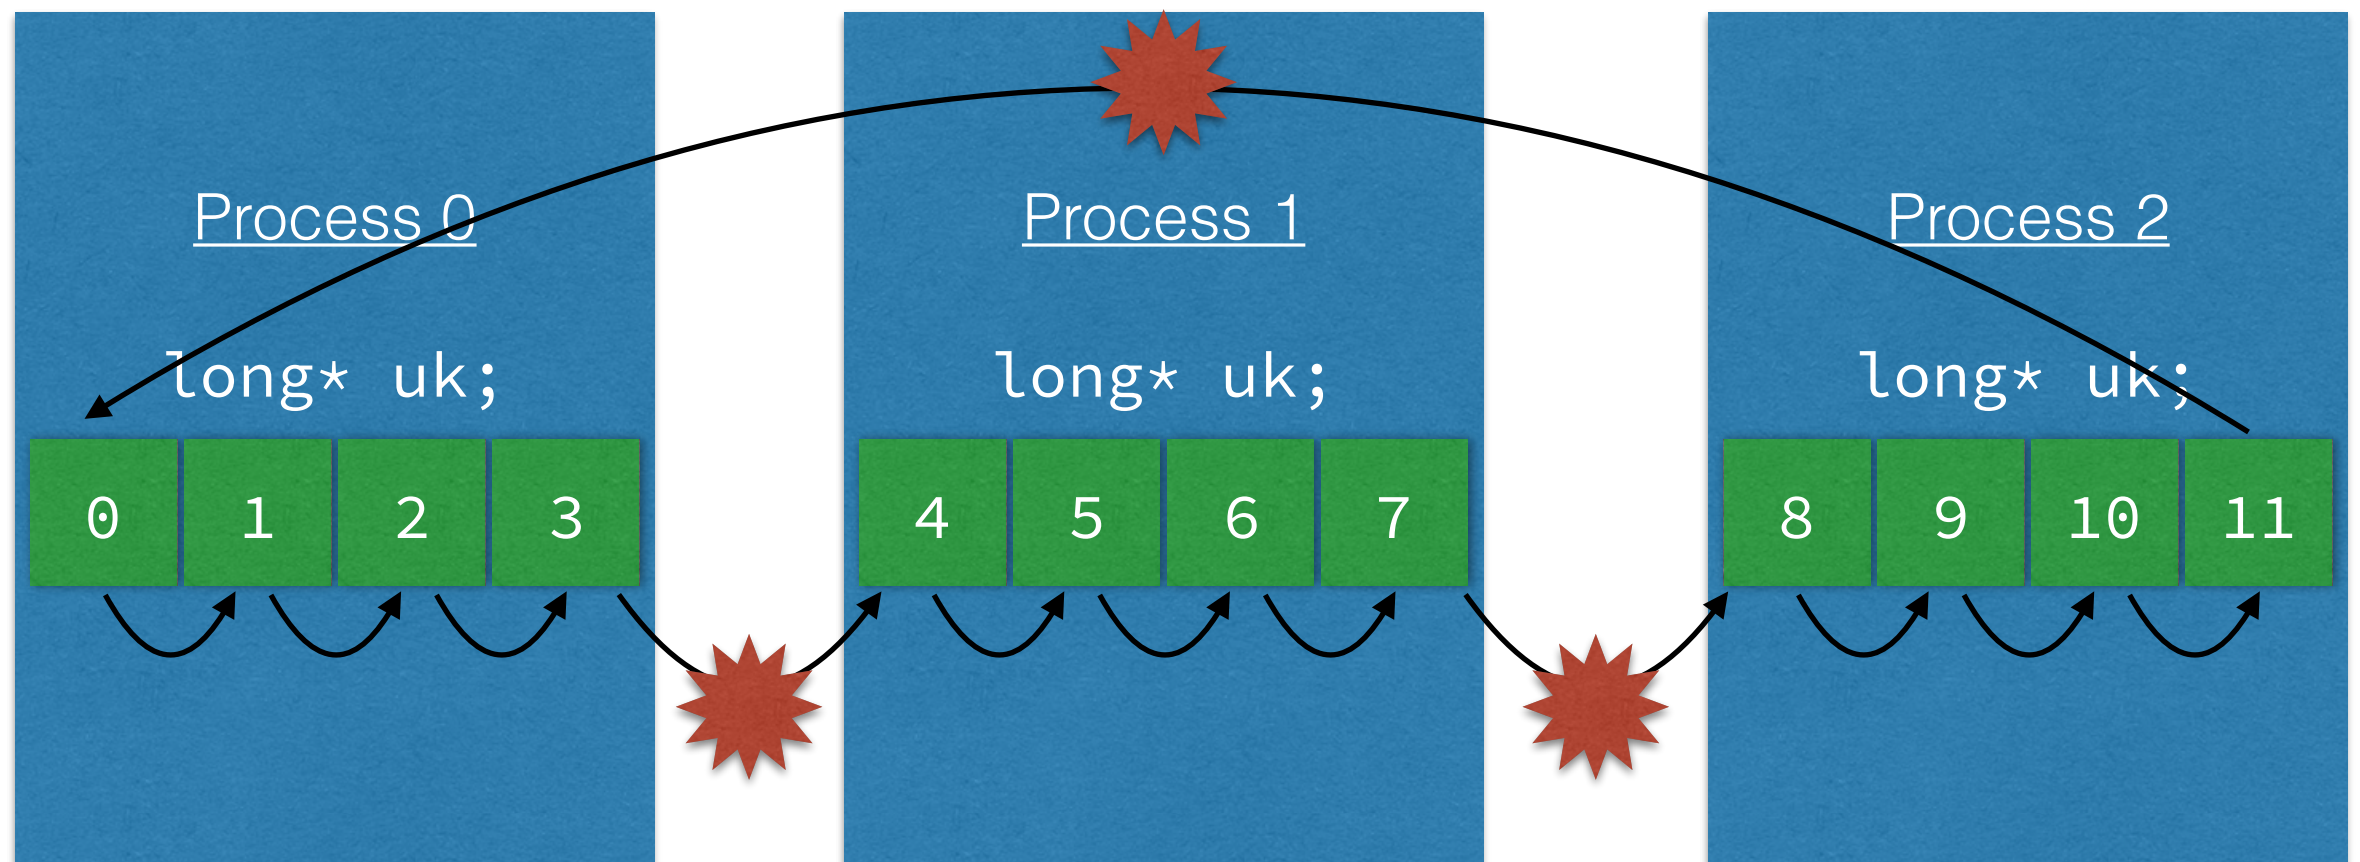
Problem Statement

- Each process allocates an array of data
- Want to “shift” the data right



Problem Statement

- Explicit Communication:
 - Process k receives data from Process $k-1$
 - Process k sends to Process $k+1$



Problem Statement


- Desired result:
- *(Optionally: shift right num_shifts times.)*



Immediate Issues

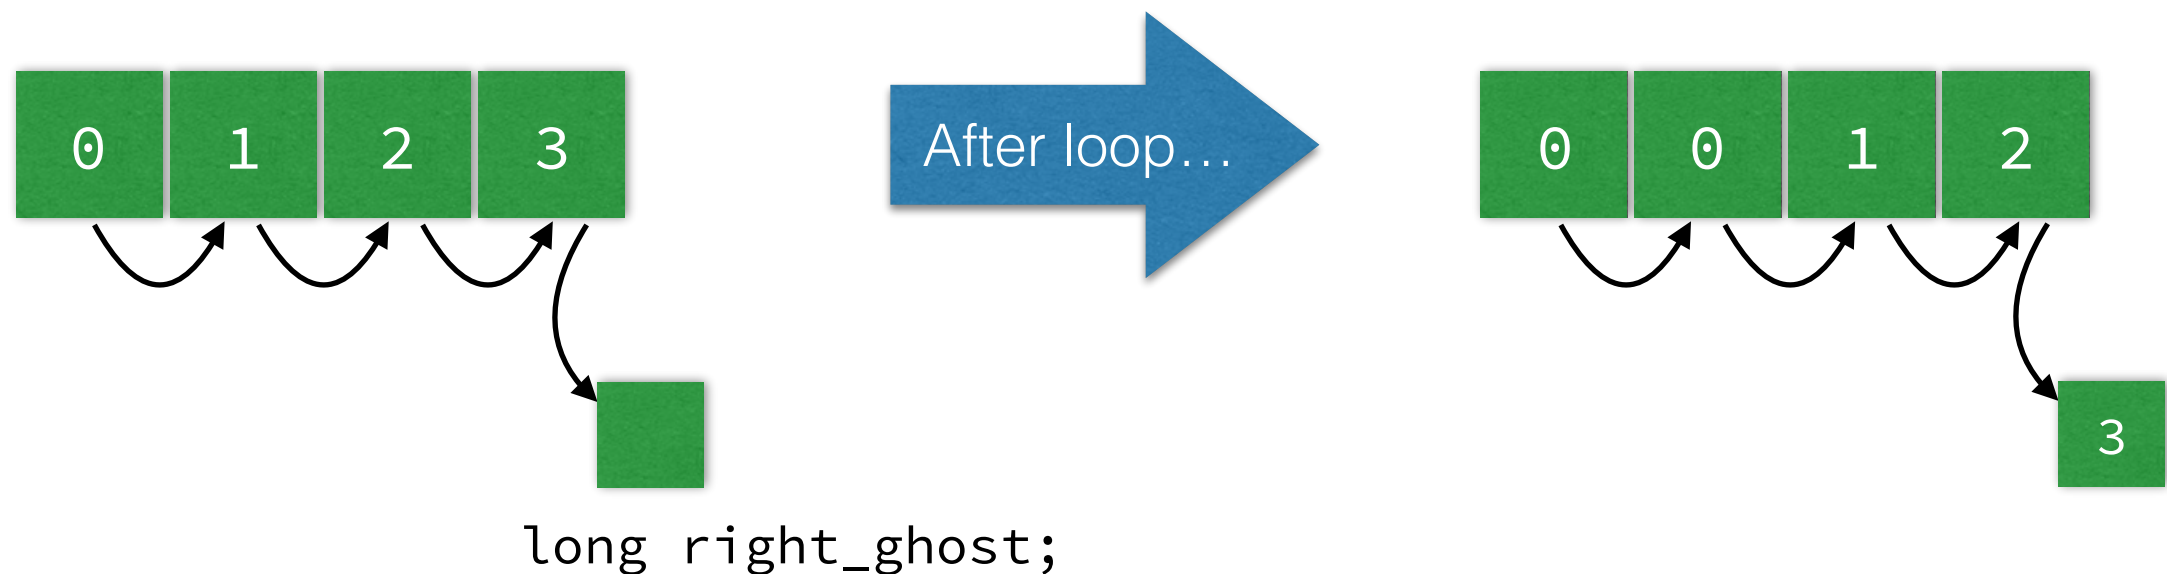
- Within each data array we risk losing information:

```
for (size_t i=N-1; i>0; --i)  
    uk[i] = uk[i-1];
```



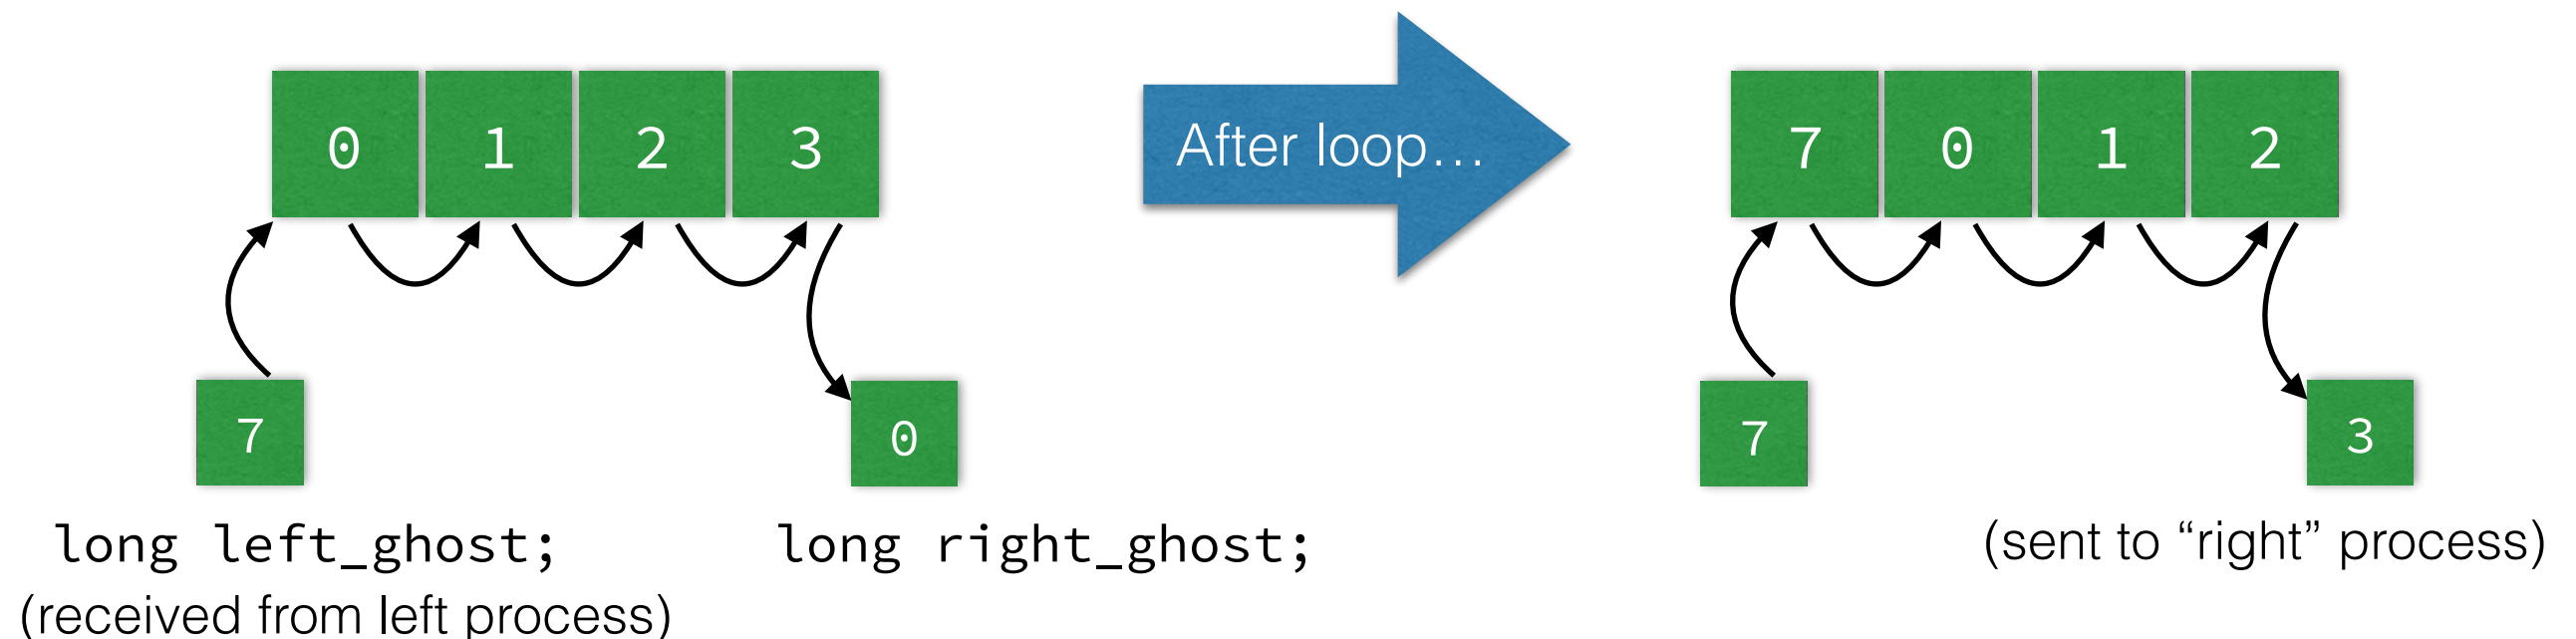
Ghost Cells

- Common domain decomposition strategy: create “ghost cells”
 - temp storage
 - data to be communicated



Ghost Cells

- Looking ahead: data incoming from “left” process
 - temp store result in `left_ghost`
 - local “comms”**: uk **global comms**: ghost cells



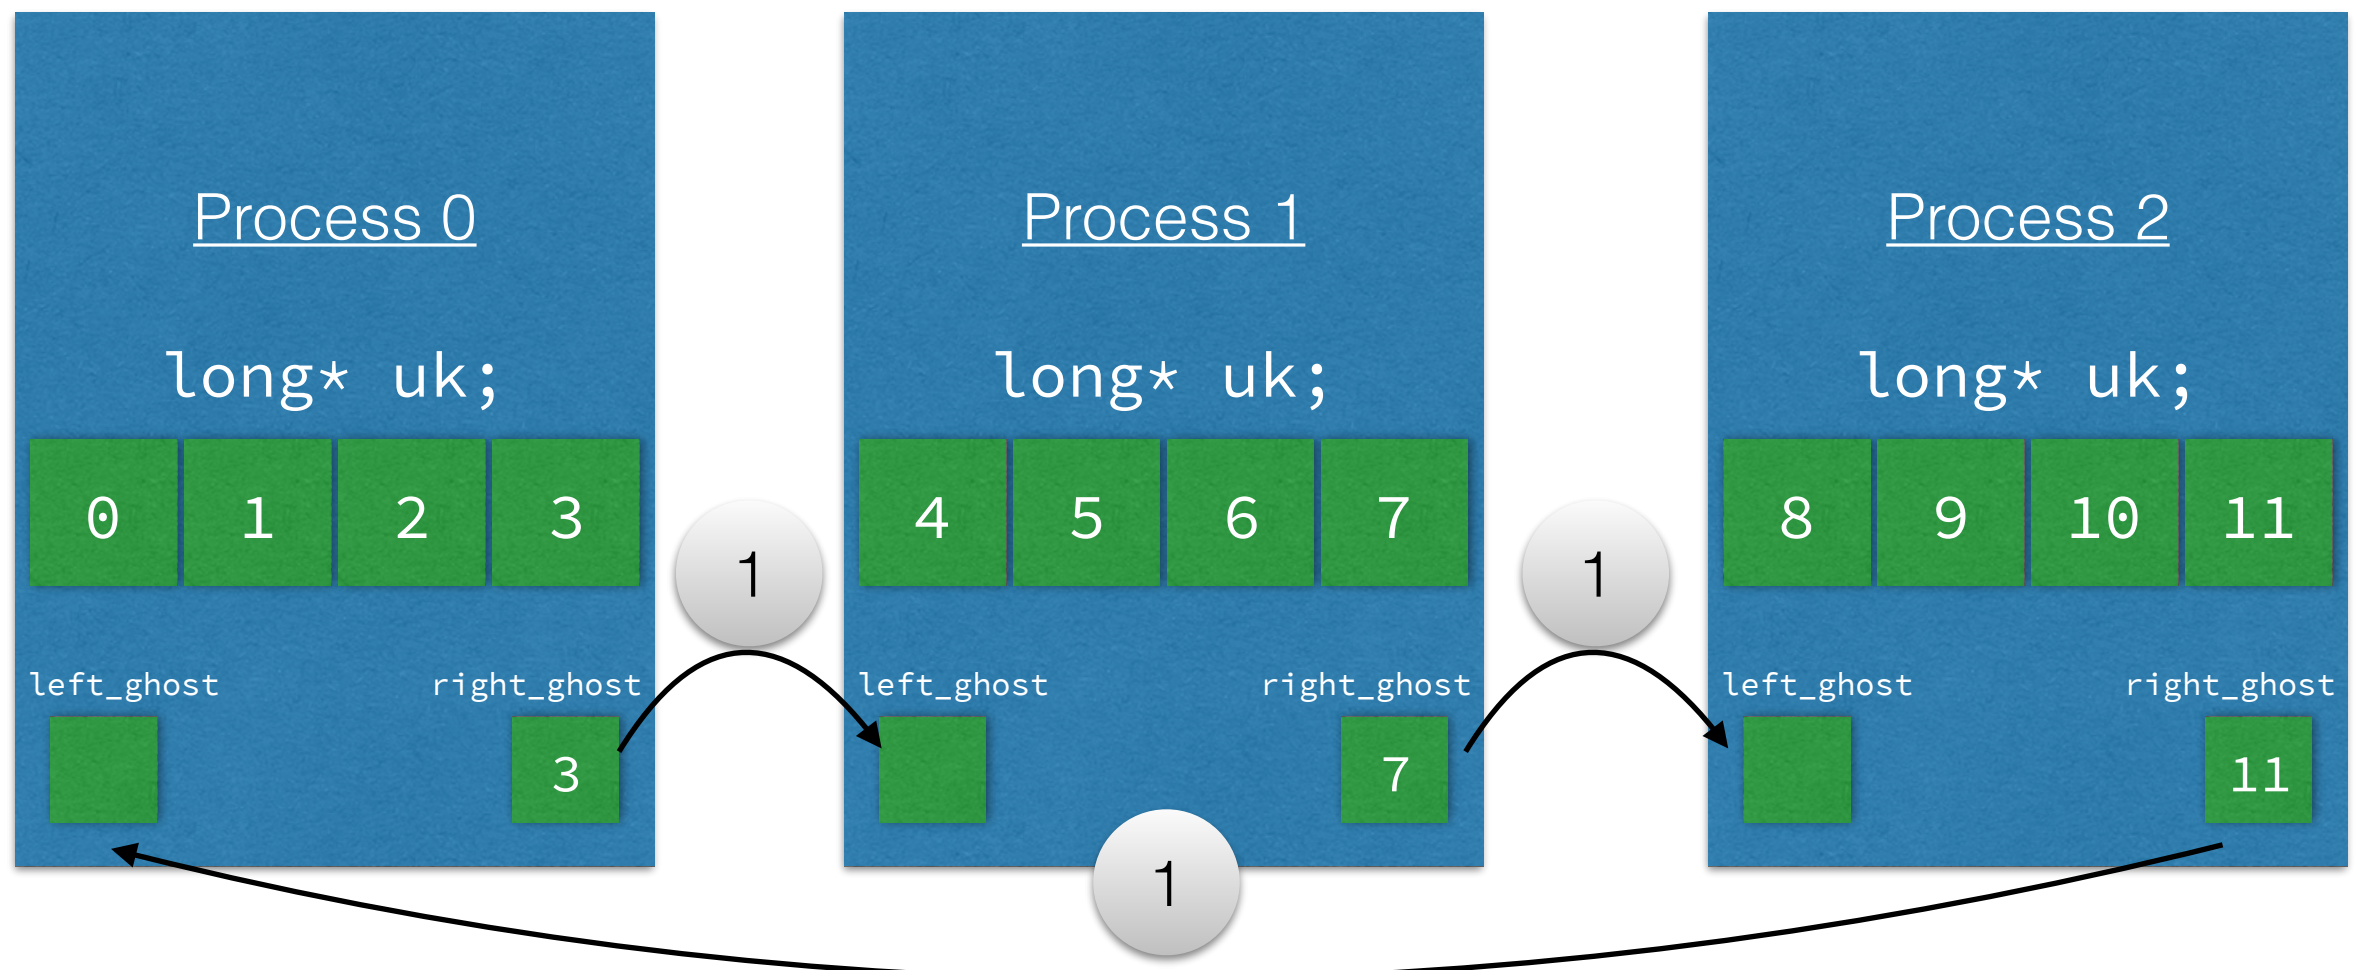
Revisiting Problem Statement

- Rewrite using ghost cells:



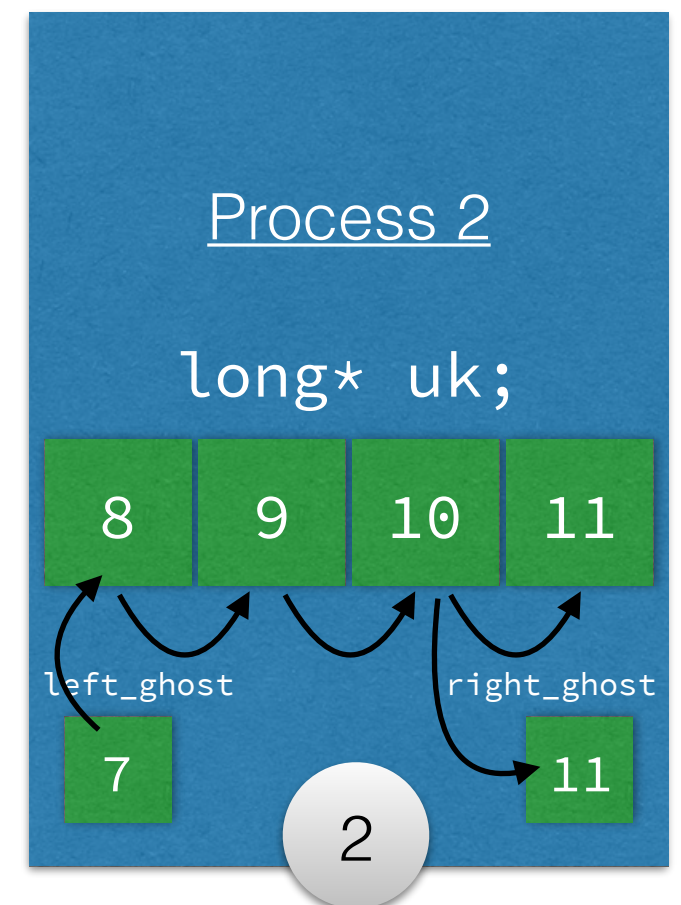
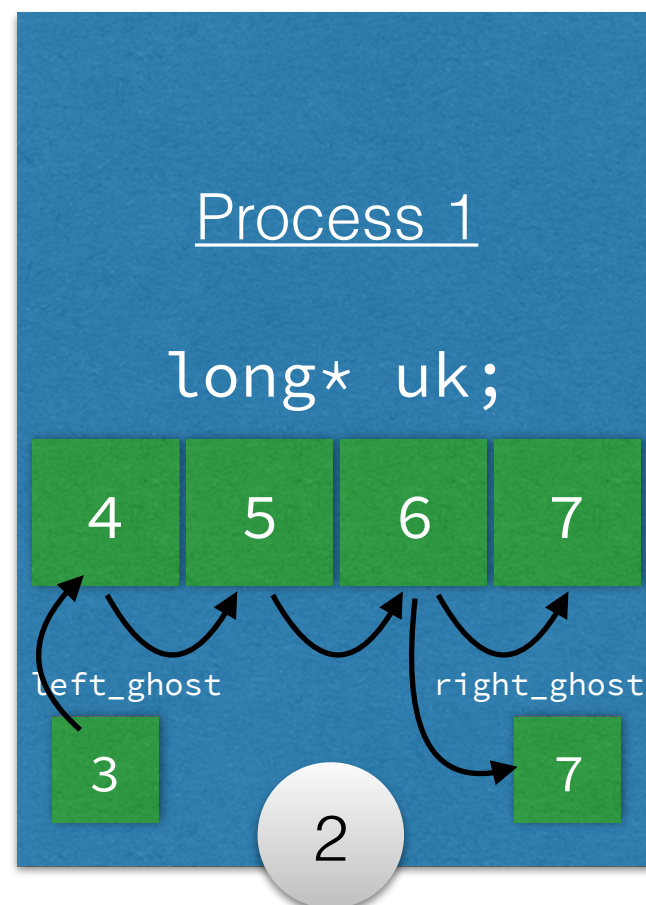
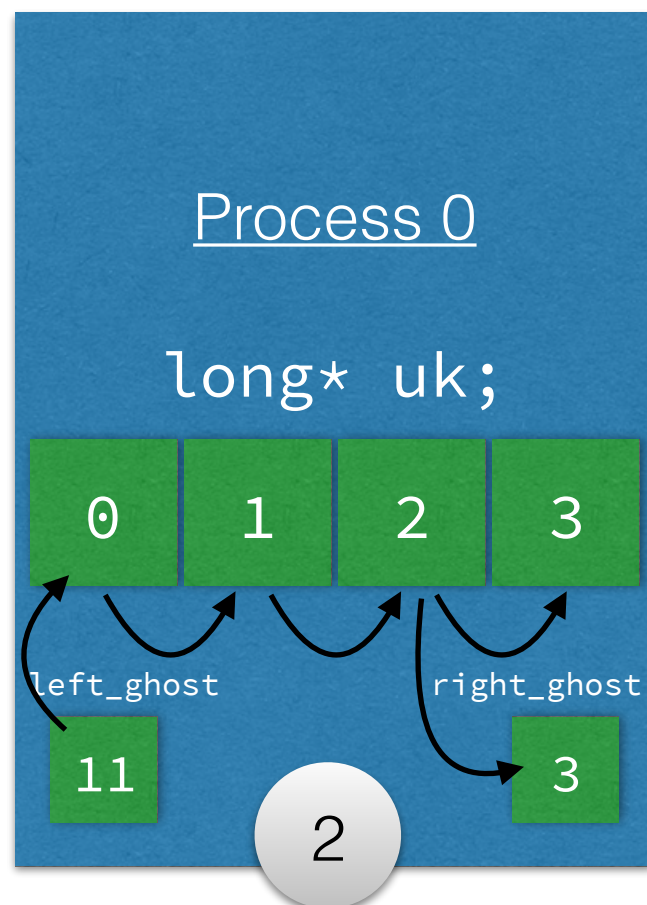
Revisiting Problem Statement

- 1) Communicate boundary data.
- 2) Perform shift within own data array.



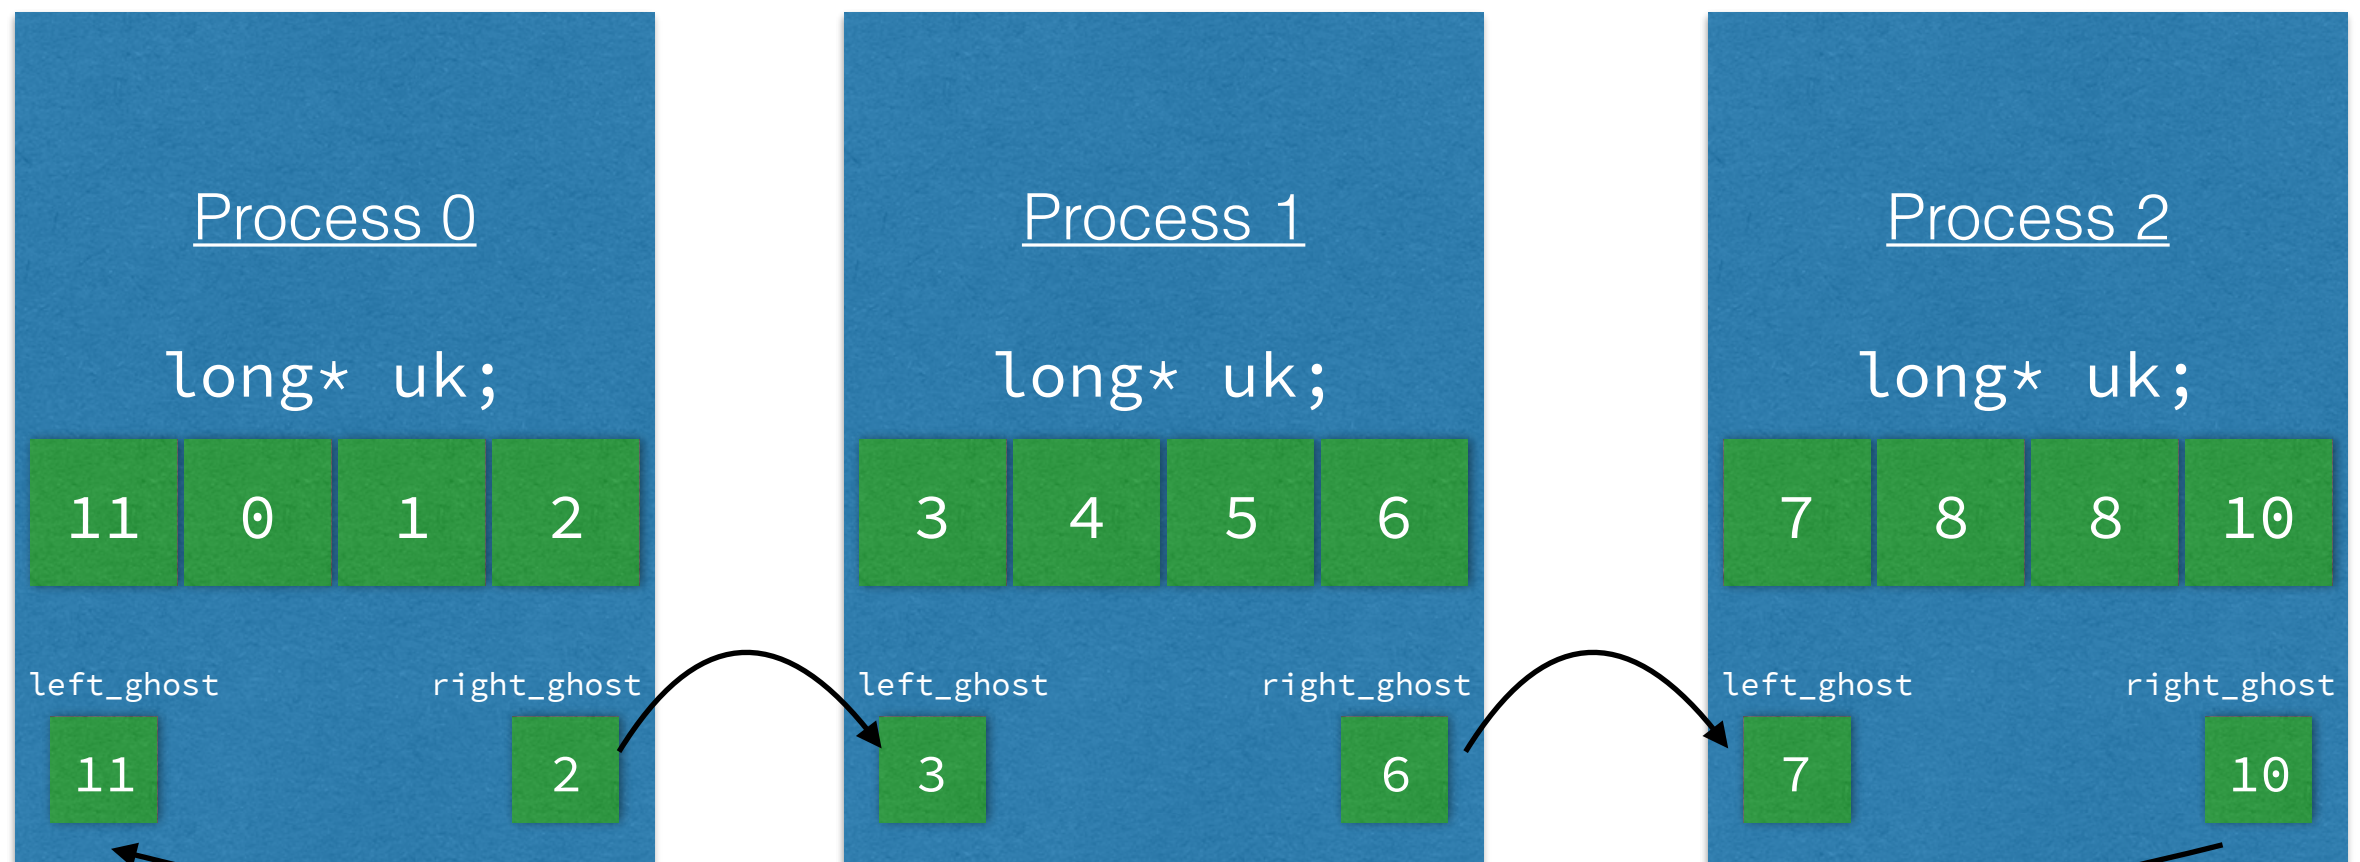
Revisiting Problem Statement

- 1) Communicate boundary data.
- 2) Perform shift within own data array.




Revisiting Problem Statement

- **Key Design Idea:** separate the local computation with the necessary communication.
- (sometimes communication latency can be hidden)



Aside: Process Topology


- Process communication laid out in ring:



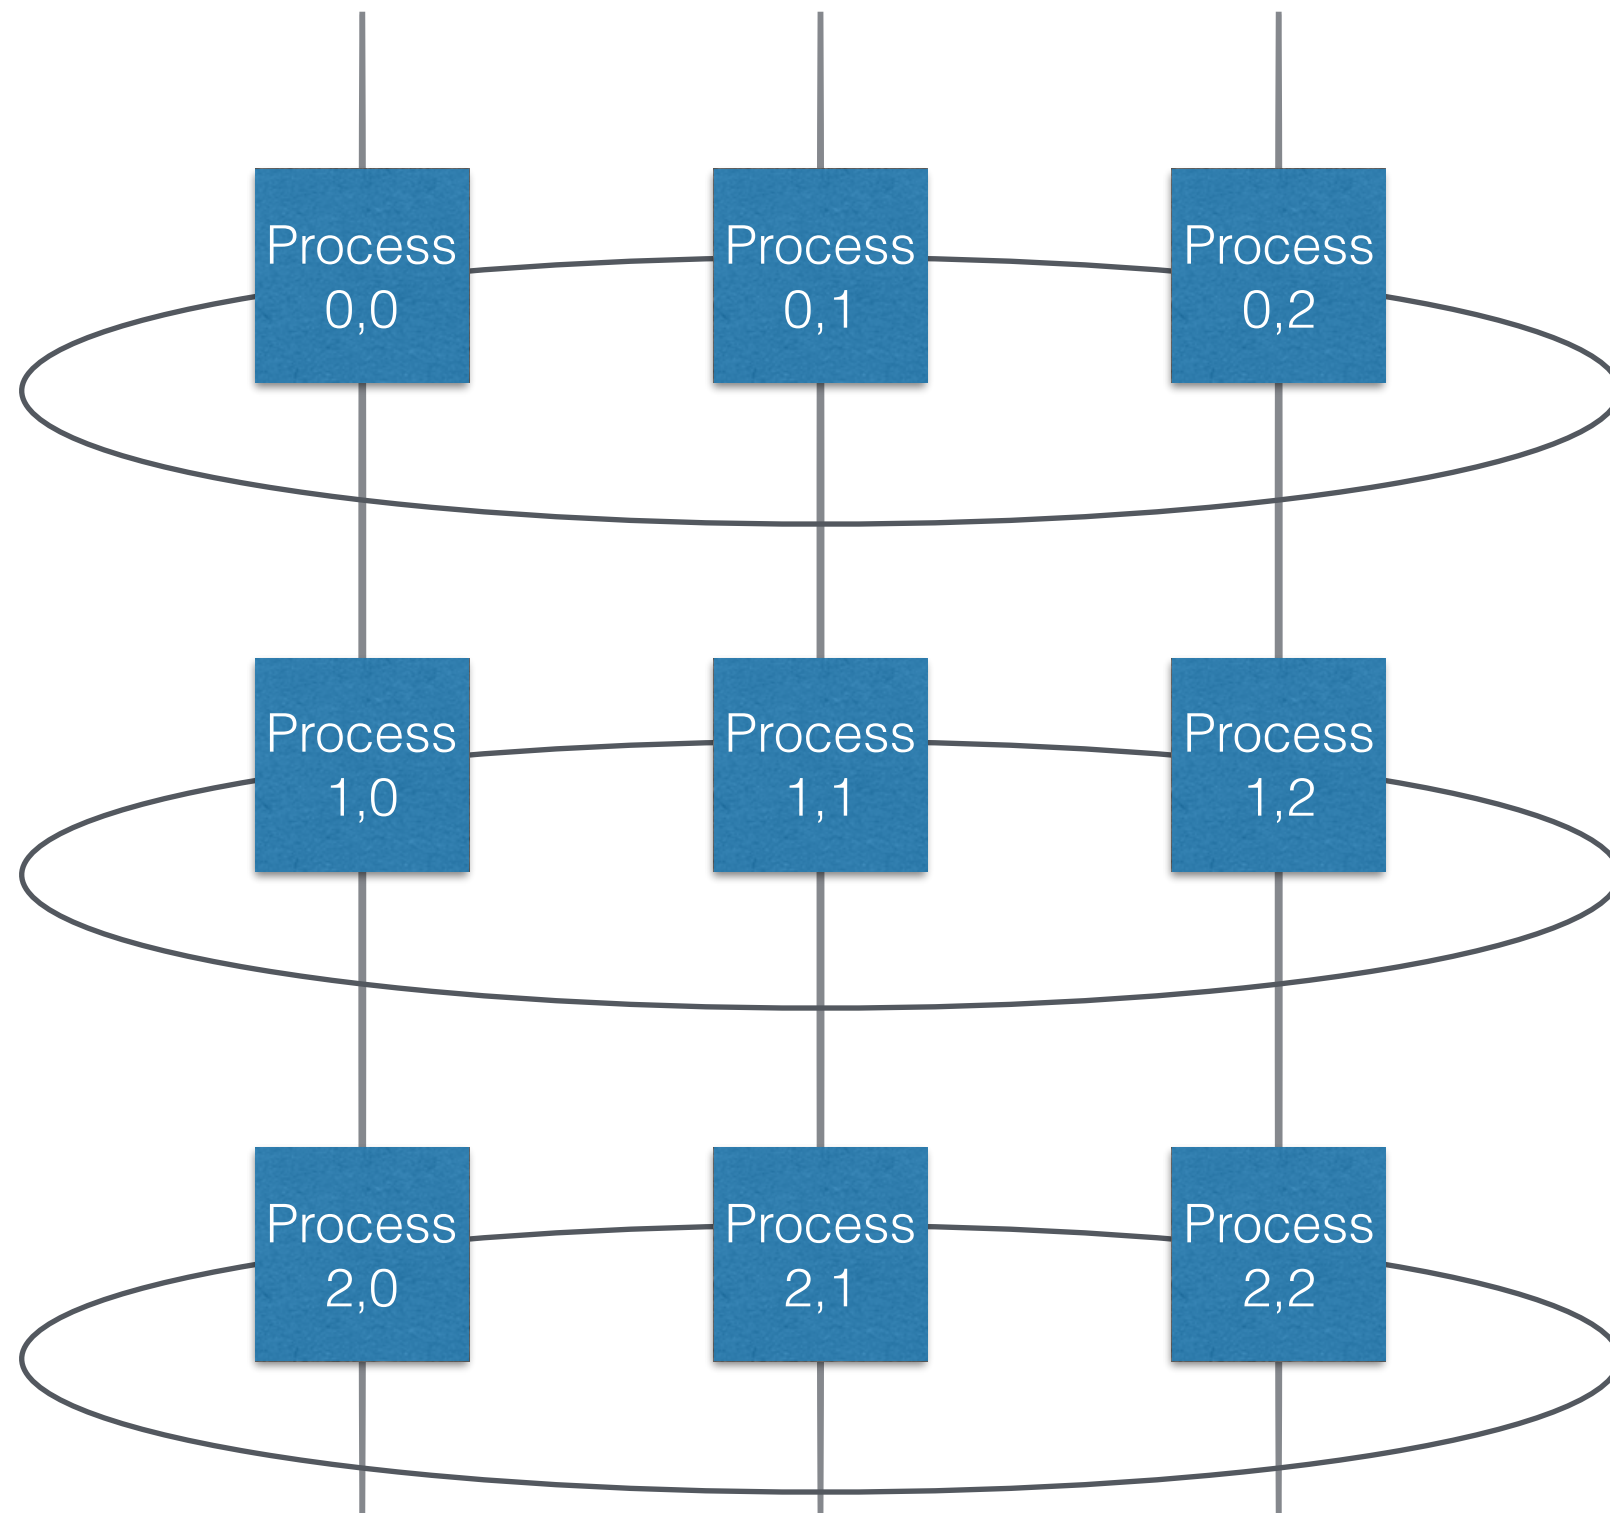
- Topology at hardware level as well. (*Certain servers connected to others in certain way.*)

Aside: Process Topology

- “Toric topology” — 2D periodic



Aside: Process Topology



Code Setup

- **Each process:**
 - gets rank and size
 - identifies “left” and “right” processes
 - allocates and initializes data, `long* uk`
 - creates ghost cell storage

Demo

Start with `shift-setup.c` and fill in “the good stuff”.
(Finished and fully documented version in `shift.c`)

Refactoring for Organization

- Suppose we want to use this `perform_shift()` functionality in another code.
 - Useful to create shared library.
 - Callable from C binaries / Python!
- Step 1: refactor within one file.

Quick Demo

`shift-refactored.c`

libshift.so

- Now, source files: `shift.h` / `shift.c`

```
void perform_shift(  
    long* uk, size_t N,  
    long num_shifts, MPI_Comm comm);
```

- Compile as usual, but with `mpicc`

```
$ mpicc -shared -o libshift.so shift.c
```

main.c

- Program: main.c

```
#include "shift.h"
int main(void) {
    // MPI setup, create data
    perform_shift(data, length, 2, MPI_COMM_WORLD)
    // MPI teardown
}
```

Number of shifts

- Compile and link:

```
$ mpicc -lshift -L. -I. main.c
$ mpiexec -n 2 ./a.out
```

Quick Demo

Contents of: `shift-library/`

mpi4py

- Miraculously, you can run MPI library code from Python using ctypes
- **Package:** mpi4py

```
$ pip install mpi4py
```

(Install may appear to freeze. Actually compiling in the background.)

- Homework #4. *(I will write wrappers for you.)*

mpi4py

- Workflow:
 1. Compile MPI C library code
 2. Link Python code via ctypes
 3. Create an mpi4py communicator object and pass to C library function
 4. Execute Python script using mpiexec:

```
$ mpiexec -n 4 python my_script.py
```


Demo

Code in `lecture15/mpi4py/` directory.